

Parental Altruism, the Samaritan’s Dilemma, and the College Decision*

Agustín Díaz Casanueva[†]

June, 2026

Latest Version

PRELIMINARY DRAFT — COMMENTS WELCOME

Abstract

I study how parental altruism shapes investment in college. Altruistic parents support children when resources are low: this finances their college attainment, insures against income risk, and—because parents cannot commit to withholding it—creates a Samaritan’s dilemma in which children under-save and over-consume. College disciplines the dilemma where cash cannot, by permanently raising the child’s income. In matched parent-child data, parents of college children consume more and transfer less often. I quantify the mechanism in a continuous-time dynastic model estimated by simulated method of moments. Altruism raises attainment: shutting it down lowers college completion from 41% to 38%, with the fall concentrated among low-ability children. The inability to commit distorts attainment unevenly—lowering it by about two percentage points on aggregate, but raising it among low-ability children while lowering it by more than twenty-five points among the higher-ability of wealthy parents. The transfers that create the dilemma also insure children, so severing the relationship to remove it lowers welfare for both parent and child—a consumption-equivalent loss of about 12% for the dynasty—while only full commitment, which preserves the insurance, raises it. The welfare effect of free college turns on the same distinction: it reallocates resources to the child (a consumption-equivalent gain of about 5%) and leaves the dynasty roughly unchanged, rather than severing the relationship that insures it.

JEL: D15, D64, I22, I23

Keywords: Parental altruism, college attendance, Samaritan’s dilemma, intergenerational transfers, family insurance, moral hazard, dynastic models

*Previously circulated as “The Role of Parental Altruism in Parents Consumption, College Financial Support, and Outcomes in Higher Education.” I am grateful to my committee members Dirk Krueger, José Víctor Ríos Rull, and Andrew Shephard for their invaluable guidance. I also thank seminar participants at the University of Pennsylvania, Universidad Católica Alumni Workshop, the Midwest Macro Meetings, and the Chilean Economics Society for helpful comments. The views expressed herein are solely those of the author and do not necessarily reflect the views of the Central Bank of Chile. All errors are my own.

[†]Central Bank of Chile, adiaz@bcentral.cl, Address: Agustinas 1180, Santiago, Chile.

1 Introduction

Parents in the United States transfer substantial resources to their adult children. Inter-vivos transfers—cash gifts, tuition, housing assistance—flow downward and are compensatory, going disproportionately to children when their resources are low and responding to their income shocks over the life cycle (McGarry, 1999, 2016)—the pattern expected when altruism, rather than exchange, is the operative motive for parent-to-child support. These transfers are most consequential around higher education, where families remain the primary source of college financing and the support a child anticipates is itself a key determinant of enrollment (Bellely and Lochner, 2007; Brown et al., 2012; Abbott et al., 2019). Two features make the setting distinctive. The support is governed by a dynamic interaction in which parents cannot commit: how much a child receives is decided as needs arise, not promised in advance. And the same anticipated support shapes two decisions at once—how much the child saves and whether the child invests in college. Family transfers and human-capital investment are therefore jointly determined within a non-committal parent–child relationship. This paper studies that interaction: how the strategic relationship between altruistic parents and their children shapes investment in college, and what it implies for the welfare effects of education policy.

At the center of the interaction is parental altruism, and its support does three things at once. It *finances* investment, supporting children at the margin of affordability through college; it *insures*, responding to the child’s realized resources and smoothing income risk; and, because parents cannot commit to withholding it, it *distorts*. The distortion takes the form of a Samaritan’s dilemma (Lindbeck and Weibull, 1988; Bruce and Waldman, 1990): children who anticipate transfers under-save and over-consume, knowing parents will step in when resources are low, inducing larger and more frequent transfers than under commitment—parents support their children *because* they care, but the anticipation of support distorts the very decisions that determine how much will be needed. Because the support is tied to the college decision, the distortion extends to human-capital investment. Two parameters therefore govern how altruism shapes college investment: its *strength*, which sets how much parents finance enrollment, and the *inability to commit*, which determines how the anticipation of support distorts it. I disentangle the two and trace each across the ability–wealth

distribution.

College plays a distinctive role in this dynamic. From the parent’s perspective, a dollar of tuition differs fundamentally from a dollar transferred: tuition permanently shifts the child’s income level, moving the child away from the states where support would be triggered, whereas a direct transfer is temporary and may itself reward under-saving. Investing in human capital therefore shrinks the *transfer region*—the set of states in which the parent would optimally provide support—disciplining the Samaritan’s dilemma through a channel cash transfers cannot replicate. The net effect on attendance is theoretically ambiguous: anticipated in-college transfers subsidize enrollment, while anticipated lifetime transfers insure children who forgo college. The structural model quantifies these opposing forces and shows that the sign of the moral-hazard effect on college varies across the ability–wealth distribution.

I formalize the mechanism in a continuous-time, non-cooperative dynastic model building on [Barczyk and Kredler \(2014a,b\)](#). Parents choose consumption and a non-negative transfer flow; children choose consumption and, at the start of the dynasty, whether to attend college. Each agent optimizes taking the other’s policy as given, while the parent internalizes the child’s welfare through an altruism parameter. In equilibrium the state space partitions endogenously into a transfer region, where the parent supports the child, and a no-transfer region, where the child is self-sufficient; the boundary—and hence the insurance value of support—is shaped by the child’s education, income, and wealth. College contracts the transfer region by permanently raising expected earnings, so the college decision depends not only on the human capital return but on how education reshapes the future financial relationship with the parent.

I document two stylized facts from the Panel Study of Income Dynamics (PSID) and the Health and Retirement Study (HRS), comparing parents of college and non-college children while conditioning on parent income. *First*, parent consumption is markedly higher when the child attends college—by about 8%, roughly \$2,500 a year. *Second*, college lowers the probability of any within-family transfer by about a fifth, while conditional amounts are weakly larger. Both line up with the central mechanism: the consumption gain far exceeds the implied fall in realized transfers, consistent with parents releasing resources because they

no longer need to *stand ready* to support the child rather than because they pay fewer transfers. The estimated model quantifies this option-value channel directly.

I estimate the model by the simulated method of moments, targeting college attainment by ability and parent wealth, the college premium, transfer patterns, and intergenerational education transmission. I first isolate the level effect of altruism: shutting it down lowers college attainment from 41% to 38%, with the fall concentrated among low-ability children—the lowest-ability quartile drops by nearly thirty percent, from 24% to 17%—for whom parental support is pivotal to enrolling at all, so altruism finances enrollment at the margin of affordability. I then decompose the enrollment effect of the lack of commitment relative to an efficient benchmark in which the parent commits to education-contingent transfers at the start of the dynasty, separating the enrollment subsidy of anticipated in-college support from the insurance that anticipated lifetime support extends to children who forgo college. On net the lack of commitment lowers aggregate enrollment modestly—by about two percentage points relative to a one-time gift—but the sign of the distortion varies across the joint distribution of ability and parental wealth: the in-college subsidy raises attendance for low-ability children (by roughly thirteen to eighteen points), while the lifetime insurance lowers it for the medium- and high-ability (by as much as twenty-seven points for medium-ability children of the wealthiest parents). The central welfare lesson concerns what is lost when the dilemma is disciplined. The transfers that distort the college margin also insure children against income risk, so severing the dynamic interaction—as a one-time transfer or an education-contingent menu would—leaves both parent and child worse off, a dynastic consumption-equivalent loss of about 12%; only full commitment, which removes the moral hazard while preserving insurance, raises welfare (by about 15% for the dynasty), and because it does so by reallocating from child (−48%) to parent (+31%) it is not a Pareto improvement. None of these commitment economies is a policy the altruistic parent can adopt—she cannot commit to withholding help—so they are benchmarks that value the dynamic interaction, not alternatives to it. The welfare effects of a free-college policy turn on the same distinction: because it leaves the relationship in place, the policy reallocates resources to the child—a consumption-equivalent gain of about 5%—while leaving the dynasty roughly neutral, rather than financing many new entrants. Studying the parent–child interaction, rather than the college decision in isolation,

is what makes these incidence and welfare results visible.

This paper connects several literatures. The first studies how families insure members against income risk. Since [Becker and Tomes \(1979, 1986\)](#), this literature has debated what motivates such transfers, distinguishing altruism, in which parents give because they value the child’s welfare, from exchange, in which transfers buy services or attention in return ([Cox, 1987](#); [Bernheim et al., 1985](#)). The evidence is mixed, and no single motive dominates across all transfer types: tests of *pure* altruism are rejected—redistributing a dollar from child to parent raises the parent’s transfer by far less than the dollar that altruism predicts ([Altonji et al., 1997](#))—and some transfers rise with recipient income, as exchange implies ([Cox, 1987](#)). For the transfers central to this paper, however—inter-vivos support flowing from parents to adult children—the altruistic margin is the operative one: such transfers are compensatory, going disproportionately to lower-income children and responding to their income shocks ([McGarry, 1999, 2016](#)). I therefore model altruism as the motive for parent-to-child support, while remaining agnostic about bequests and old-age assistance, where exchange and accidental components loom larger. How complete the resulting insurance is remains debated: [Altonji et al. \(1992\)](#) and [Hayashi et al. \(1996\)](#) reject perfect family insurance, yet families still insure substantially—[Attanasio et al. \(2018\)](#) finds considerable insurance potential between parents and adult children, and [Kaplan \(2012\)](#) shows the option to move home insures young workers. Closest, [Boar \(2021\)](#) identifies dynastic precautionary saving; I build on this, showing college attenuates the dynastic precautionary motive by raising the child’s income and shrinking the transfer region. Transfers also fund large, lumpy investments and transmit wealth through gifts and bequests ([Nishiyama, 2002](#)): [Barczyk et al. \(2023\)](#) emphasize how illiquid housing delays transfers until death, and [Brandsaas \(2025\)](#) finds parental transfers account for a sizable share of young-adult homeownership by relaxing down-payment constraints. I study the same insight—that family transfers shape children’s long-run investment—for human capital.

A second strand concerns the trade-off between commitment and the insurance that an ongoing, non-committal relationship provides, and my welfare results sit squarely within it. A long tradition shows that informal, self-enforcing insurance arrangements sustain consumption smoothing precisely because they remain flexible and state-contingent, and that a formal

or committed substitute can crowd them out: [Attanasio and Ríos-Rull \(2000\)](#) show that public insurance, by improving members’ outside options, can unravel the private risk-sharing of village economies and reduce welfare. The same logic governs the family here. Replacing the parent’s ongoing, state-contingent transfers with a one-time transfer severs the insurance the relationship provides, so disciplining the Samaritan’s dilemma can leave the dynasty—parent and child alike—worse off. More broadly the tension is one of commitment versus flexibility ([Amador et al., 2006](#)): commitment disciplines a behavioral distortion—there, an individual’s temptation to over-consume; here, the child’s strategic under-saving—but forfeits the responsiveness to shocks that makes the arrangement valuable as insurance. My contribution is to show that it is this trade-off, rather than commitment per se, that determines the welfare value of the relationship. The altruistic parent can commit to nothing: a promise to follow a state-contingent schedule, or simply to “transfer once and stand back,” is time-inconsistent, since whenever the child nears the transfer region $\eta > 0$ makes her want to help. The commitment economies are therefore benchmarks, not policies, and they differ only in what they would preserve—a one-time transfer severs the state-contingent flexibility that is the relationship’s insurance value, whereas full commitment would keep it. This is the sense in which the result departs from the canonical limited-commitment risk-sharing literature ([Kocherlakota, 1996](#); [Ligon et al., 2002](#)), in which greater commitment delivers greater insurance: there the binding friction is enforcement, while here it is the parent’s own altruism, and the relationship’s value lies precisely in the state-contingent flexibility that a one-time commitment would sacrifice.

Because parents here hold wealth to stand ready to support adult children, the paper connects to the literature on slow wealth decumulation in retirement. [De Nardi et al. \(2010\)](#) trace much of it to longevity and medical-expense risk and [Love et al. \(2009\)](#) show annualized wealth is roughly flat over retirement; family links do much of the rest, as [De Nardi \(2004\)](#), [De Nardi and Yang \(2014\)](#), and [Lockwood \(2018\)](#) show voluntary bequests reproduce the slow drawdown and the concentration of wealth, and [Ameriks et al. \(2016\)](#) find parents hold wealth to assist children when need is greatest. This literature models the family motive largely as a terminal bequest. Building on the forward-looking inter-vivos motive of [Barczyk and Kredler \(2014a,b\)](#), I instead let parents retain wealth for the option value of an endogenous transfer

region whose extent depends on the child’s college, and I target the pace of late-life drawdown directly. A related literature asks whether financial constraints bind for college: [Cameron and Heckman \(1998\)](#) and [Keane and Wolpin \(2001\)](#) found a small role for family income conditional on ability, but [Belley and Lochner \(2007\)](#) document its growing importance and [Lochner and Monge-Naranjo \(2011\)](#) find binding constraints for some students, while [Brown et al. \(2012\)](#) stress that the expected family contribution is neither guaranteed nor universal and [Hotz et al. \(2018\)](#) show parental income and housing wealth affect attendance and graduation. Quantitatively, [Restuccia and Urrutia \(2004\)](#), [Caucutt and Lochner \(2020\)](#), and [Heckman and Mosso \(2014\)](#) study how parental investment, borrowing constraints, and family environments shape skills and persistence. Closest, [Abbott et al. \(2019\)](#) show government aid crowds out private contributions; I add a channel—moral-hazard reduction through the Samaritan’s dilemma—through which parental wealth affects enrollment beyond direct cost subsidies.

Methodologically, the model builds on the quantitative literature on family dynamics without commitment ([Laitner, 1988](#); [Lindbeck and Weibull, 1988](#)), in which [Lee and Se-shadri \(2019\)](#) develop a dynastic model with parental human-capital investment. I adopt the continuous-time framework of [Barczyk and Kredler \(2014a,b\)](#), which yields a well-defined equilibrium with an endogenous transfer region and has been applied to long-term care and housing ([Barczyk and Kredler, 2018](#); [Barczyk et al., 2023](#)); I extend it by embedding endogenous college decisions. The Samaritan’s dilemma at the center of the mechanism ([Lindbeck and Weibull, 1988](#); [Bruce and Waldman, 1990](#)) here takes the form of strategic under-saving by the child, which college mitigates by permanently shifting income and reducing entry into the transfer region.

The remainder of the paper is organized as follows. Section 2 presents the model. Section 3 documents the stylized facts. Section 4 describes the estimation strategy. Section 5 discusses the model fit and the role of parental altruism. Section 6 introduces the full commitment benchmark, formalizes the two opposing forces of moral hazard on college, and presents the quantitative decomposition. Section 7 conducts the welfare analysis. Section 8 concludes.

2 Model

This section presents a continuous-time, non-cooperative dynastic model without commitment that captures the interactions between altruistic parents and their adult children. The model extends the framework of [Barczyk and Kredler \(2014a,b\)](#) by incorporating endogenous college decisions, allowing me to study how family transfers affect college financial support, graduation, and parent retirement consumption. I first describe the demographic structure and income processes, then present the coupled parent-child optimization problem, the parent's transfer decision, and the college decision.

2.1 Demographics and Timing

Time is continuous. Each dynasty consists of one parent and one child who overlap for a finite period, as illustrated in [Figure 1](#). The child enters the model at age 18, at which point the parent is 42. Parent and child are coupled for the parent's entire remaining life: the overlap runs from child age 18 / parent age 42 to the parent's death at age 72, when the child is 48. The parent works until age 66 ($t = T^{\text{ret}}$), earning labor income $y_p(t)$, and then receives social security income $\text{SS}(e_p)$ until death at age 72 ($t = T$, where $T = T^{\text{ret}} + 6$). The retired parent remains coupled with the child, so transfers $\tau(t) \geq 0$ are available throughout the overlap $[0, T]$, including during the parent's retirement; the coupled problem is solved over the parent's full remaining life. Death is deterministic at age 72 (the model has no mortality risk). At the parent's death, a fraction α_{beq} of remaining parent wealth transfers to the child as a bequest; the remaining $(1 - \alpha_{\text{beq}})$ share is not transmitted.¹ The child, now 48, then continues alone—working until its own retirement at 66 and consuming social security until death at 72. The child in turn becomes the parent of a child of its own, so that its terminal state—ability, education, and wealth—becomes the starting point for the next generation, and the dynasty continues.

¹The bequest fraction α_{beq} separates the two roles that terminal wealth plays in the model. The warm-glow motive ϕ_b requires parents to hold substantial wealth at death—this is what matches the absence of decumulation at ages 60–70—while the conditional child-wealth moments (child wealth conditional on parent wealth) discipline how much of that wealth children actually receive. The untransmitted share $(1 - \alpha_{\text{beq}})$ is a reduced-form stand-in for end-of-life wealth absorption that the model's compressed lifespan omits: death is deterministic at 72, so the later years in which out-of-pocket medical and long-term-care spending absorb much of retirement wealth ([De Nardi et al., 2010](#); [Lockwood, 2018](#)) are not modeled, and the wealth that would be consumed in them is treated as leaving the dynasty rather than reaching the child.

At the beginning of the dynasty, the child makes a one-time, irreversible college decision. Children who attend college spend four years (ages 18–22) earning a fraction ℓ_C of full-time income—paid at the high-school wage, since the degree is not yet completed—while paying annual tuition ϕ . After college, the child enters the labor force with a permanently higher income process. Children who do not attend college enter the labor market immediately at age 18.


Both parent and child can save in a risk-free bond paying interest rate r . The child borrows through the student loan during and after college, paying the fixed student-loan rate $r + \iota_s$ on the balance, which is amortized to zero by age 32 (the standard ten-year repayment plan).² Beyond the student loan, both the child and the parent may hold unsecured debt as working-age adults: following [Abbott et al. \(2019\)](#), education-specific limits of \$85,000 for college graduates and \$25,000 otherwise apply, the outstanding balance carries the estimated wedge $r + \iota_b$, and the limit ramps linearly to zero over the years before retirement, so no one retires or dies in debt (a no-Ponzi constraint) and bequests remain non-negative. Markets are incomplete: no state-contingent insurance is available against the child’s income risk. The parent and child are connected through four financial links: a one-time college-conditional gift paid at enrollment, a continuous transfer flow that supports the child’s consumption during the overlap, discrete lump-sum transfers that arrive at exogenous Poisson opportunities and move wealth between generations, and a bequest at the parent’s death, when a fraction α_{beq} of remaining parent wealth transfers to the child. The flow alone cannot move wealth across generations—capped at the child’s desired consumption, it only supports spending—so the lump-sum transfers and the bequest are the channels through which parental wealth actually reaches the child: the former match the discrete, life-event-driven gifts observed in the data, the latter captures transmission at death. Because the lump opportunities arrive exogenously rather than on a schedule the parent chooses, adding them leaves the no-commitment structure intact.

Two features deserve comment. First, altruism is one-sided: the parent cares about the child’s welfare, but the child does not internalize the parent’s utility. Second, transfers flow only downward—from parent to child—and the overlap is finite, ending at the parent’s death

²The federal Standard Repayment Plan, the default plan for Direct Loans, fixes repayment at 120 equal monthly payments over ten years ([U.S. Department of Education, Federal Student Aid, 2025](#)).

(age 72).³

Figure 1. Dynasty Timeline



Notes: Two consecutive generations. The parent (upper line) and child (lower line) overlap for the parent's remaining life (parent ages 42–72, child 18–48; shaded), with the child's college phase dashed (18–22, decision $\diamond$). Four transfer channels run from parent to child: a college gift at enrollment (τ_C) and a continuous consumption-support flow ($\tau(t)$, dashed arrows); discrete lump-sum transfers at Poisson arrivals after college (τ_L , solid arrows); and a bequest at the parent's death ($\alpha_{\text{beq}} a_p$, bold arrow). The child's terminal state then initializes the next generation's parent (curved arrow), and the dynasty continues. Ages and labels are shown once, on generation n .

2.2 State Variables and Income Processes

The state of a dynasty at any instant t during the overlap phase is $(a_p(t), a_c(t), z(t); \theta, z_p, e_p, e_c)$, where a_p and a_c are parent and child wealth, z is the child's stochastic productivity shock, θ is the child's cognitive ability, z_p is the parent's persistent productivity component (initialized at the dynastic hand-off by the state the parent held as the previous generation's child, and evolving over the parent's working life), and $e_p, e_c \in \{HS, C\}$ denote the education levels of parent and child.

³McGarry (1999) documents that inter-vivos transfers from parents to children are substantial and go disproportionately to less well-off children, and McGarry (2016) shows that they respond to children's incomes over the life cycle; transfers from adult children to parents are comparatively rare, as old-age consumption is financed primarily through Social Security, pensions, and private savings.

Child income. The child’s labor income at time t is:

$$y_c(t) = w_{ec} \cdot \left(\frac{\theta}{\bar{\theta}}\right)^{\lambda_{ec}} \cdot \exp(\gamma(t, e_c) + z(t))$$

where $h(\theta, e_c) = (\theta/\bar{\theta})^{\lambda_{ec}}$ maps ability into human capital with an education-specific ability gradient λ_{ec} , following [Abbott et al. \(2019\)](#), and $\bar{\theta} = \exp(\mu_\theta)$ is the population mean ability.⁴ The gradient $\lambda_C > \lambda_{HS}$ implies that cognitive ability and college education are complements: higher-ability individuals benefit more from college. The education-specific wage w_{ec} captures the relative price of education-specific labor: $w_{HS} = w$ and $w_C = w \cdot \omega_C$. The deterministic age-income profile is quadratic in age, $\gamma(t, e_c) = \gamma_{1,e} t - \gamma_{2,e} t^2$, and $z(t)$ follows an Ornstein-Uhlenbeck process:

$$dz = -\kappa_{ec} z dt + \sigma_{ec} dW$$

with education-specific mean-reversion rate κ_{ec} and diffusion σ_{ec} ([Appendix D](#)). The channel through which college reshapes the child’s income distribution, and hence the transfer region, is the income *level*: a higher ability gradient ($\lambda_C > \lambda_{HS}$), a steeper age-income profile ($\gamma_{1,C} > \gamma_{1,HS}$), and the relative wage ω_C jointly raise expected income and keep college-educated children far from the states where parental transfers would be triggered. During college ($t \in [0, 4]$ for $e_c = C$), the child works part-time, earning a fraction ℓ_C of full-time earnings. Because the degree has not yet been completed, these earnings are paid at the high-school level—the wage w_{HS} , ability gradient λ_{HS} , and age profile $\gamma(t, HS)$ —and the college premium applies only after graduation.⁵

Parent income. Parent income depends on education, age, and a persistent productivity component $z_p(t)$ that evolves over the parent’s working life: before retirement, $y_p(t) = w_{ep} \cdot (\theta/\bar{\theta})^{\lambda_{ep}} \cdot \exp(\gamma(t, e_p) + z_p(t))$, where z_p follows the same education-specific Ornstein-Uhlenbeck process as any worker, $dz_p = -\kappa_{ep} z_p dt + \sigma_{ep} dW_p$, initialized at the dynastic hand-off (age 42) by the productivity state the parent held as the previous generation’s child; after retirement at age 66, $y_p(t) = SS(e_p)$. The parent therefore faces genuine, persistent

⁴The normalization by $\bar{\theta}$ ensures that the human capital component equals one at mean ability for both education groups, so that the average college premium comes entirely from the life-cycle profile and wage price ω_C , not from the level effect of $\lambda_C \neq \lambda_{HS}$ evaluated at a non-zero mean $\log \theta$.

⁵In-college earnings therefore rise only weakly with ability, so high-ability students cannot self-finance net tuition out of part-time work and parental resources remain relevant throughout the ability distribution.

labor-income risk over the overlap. This risk is a second source of within-dynasty uncertainty alongside the child’s income shock z : it generates a precautionary saving motive and the cross-sectional dispersion of parent income and wealth, and—through the transmission channel below—a direct dependence of the child’s earning capacity on the parent’s realized income.⁶

Ability transmission. Cognitive ability is transmitted across generations through a log-normal AR(1) process: $\log \theta' = \rho_\theta \log \theta + \mu_\theta(1 - \rho_\theta) + \epsilon_\theta$, with $\epsilon_\theta \sim N(0, \sigma_\theta^2)$ and μ_θ the unconditional mean of log ability. The drift term $\mu_\theta(1 - \rho_\theta)$ ensures that the stationary distribution of ability has mean μ_θ , generating realistic intergenerational persistence while allowing for regression to the mean.

Intergenerational earnings transmission. Beyond ability and bequests, a third link operates through initial earning capacity. A child enters the labor market with an initial productivity draw $z_0 \sim N(\zeta_{pe} \mathbf{1}\{e_p = C\} + \rho_{zp} z_p^{\text{entry}}, \sigma_{z_0, e_c}^2)$, where $\zeta_{pe} \geq 0$ is a parent-education earnings premium, $\rho_{zp} \geq 0$ an intergenerational income-transmission coefficient (letting earning capacity depend on the parent’s income *level*, not only education; [Lee and Seshadri, 2019](#)), and z_p^{entry} the parent’s persistent productivity at the child’s entry. Because the model begins at age 18, the childhood through which parental background shapes earning capacity ([Restuccia and Urrutia, 2004](#); [Lee and Seshadri, 2019](#)) is pre-determined; ζ_{pe} and ρ_{zp} are its reduced form, the counterpart of the parental dependence of skill endowments in [Abbott et al. \(2019\)](#). Since the early-career advantage accumulates permanently into wealth even as the mean-reverting shock z decays, this channel routes part of the parent–child wealth correlation through inherited earning capacity rather than liquid transfers, leaving the child’s age-18 college-financing constraint untouched; absent it, the wealth-transmission moments would load entirely on ability and the financial channels, overstating the role of family resources in college attendance.

⁶The process is initialized at age 42 by the persistent state the parent held as the previous generation’s child and continues to evolve thereafter, generating the parent-income dispersion endogenously.

Asset-return risk. Following [Barczyk and Kredler \(2014a\)](#), each agent’s wealth carries a small multiplicative shock,

$$da_p = \dot{a}_p dt + \sigma_a a_p dB_p, \quad da_c = \dot{a}_c dt + \sigma_a a_c dB_c,$$

where B_p and B_c are independent standard Brownian motions, the drifts $\dot{a}_p, \dot{a}_c$ are given in (8)–(9), and the volatility, $\sigma_a \max(a, 0)$, vanishes at zero wealth and is absent on debt. The shock represents idiosyncratic asset-return risk and, following [Barczyk and Kredler \(2014a\)](#), regularizes the equilibrium: a diffusion in the *wealth* state—which the income shock z cannot provide—renders the value functions C^2 and the transfer boundary smooth, restoring a pure-strategy Markov equilibrium and eliminating the “blast from the past” of [Barczyk and Kredler \(2021\)](#) (Appendix F). I set σ_a to about ten percent per year, a realistic diversified-portfolio return standard deviation, at which it both regularizes the equilibrium and, as genuine return risk, contributes to the cross-sectional dispersion of wealth.

Return heterogeneity. On positive balances the parent earns a return that rises with the level of wealth,

$$r_{\text{eff}}(a_p) = r + r_{a,\text{grad}} \frac{a_p}{a_p + a_{\text{ref}}},$$

a bounded, saturating premium that vanishes near the median ($a_p \ll a_{\text{ref}}$) and approaches $r + r_{a,\text{grad}}$ for the wealthy. This scale dependence in returns—documented in the wealth-inequality literature ([Fagereng et al., 2020](#); [Gabaix et al., 2016](#))—concentrates the right tail of the wealth distribution, lifting the p90/p50 ratio and aggregate wealth through the top without shifting the median or the pace of decumulation; and because the parent’s productivity type is inherited, it also contributes to intergenerational wealth persistence. The premium applies only to positive assets, with the saturation scale a_{ref} fixed and the slope $r_{a,\text{grad}}$ estimated; debt instead carries the borrowing wedges defined below. Setting $r_{a,\text{grad}} = 0$ recovers a homogeneous return.

2.3 The Coupled Parent-Child Problem

During the overlap period, parent and child simultaneously choose consumption and savings, while the parent additionally chooses a non-negative transfer flow $\tau(t) \geq 0$. The interaction

is non-cooperative: each agent optimizes taking the other's policy as given, but the parent internalizes the child's welfare through the altruism parameter $\eta > 0$.

The continuous-time formulation follows [Barczyk and Kredler \(2014a\)](#), who show that simultaneity of moves in continuous time yields a well-defined equilibrium—an advantage over discrete-time formulations where the timing protocol (who moves first within a period) can affect the equilibrium outcome.

2.3.1 Parent's Problem

The parent chooses consumption c_p and transfer flow $\tau \geq 0$ to maximize lifetime utility, weighting the child's instantaneous utility by η . The parent's value function $V^p(a_p, a_c, z, z_p; t)$ —which depends on the parent's own persistent productivity state z_p as well as the child's—satisfies:

$$\begin{aligned} \rho V^p = \max_{c_p, \tau \geq 0} & \left\{ u(c_p) + \eta u(c_c^*) + V_{a_p}^p [r_{\text{eff}}(a_p) a_p + y_p - c_p - \tau] + V_{a_c}^p [\tilde{r}_c a_c + y_c - c_c^* + \tau] \right. \\ & + V_z^p(-\kappa z) + \frac{1}{2} \sigma^2 V_{zz}^p + V_{z_p}^p(-\kappa_p z_p) + \frac{1}{2} \sigma_p^2 V_{z_p z_p}^p + \frac{1}{2} \sigma_a^2 a_p^2 V_{a_p a_p}^p + \frac{1}{2} \sigma_a^2 a_c^2 V_{a_c a_c}^p + V_t^p \\ & \left. + \lambda_g [V^p(a_p - \tau_L^*, a_c + \tau_L^*, z, z_p; t) - V^p(a_p, a_c, z, z_p; t)] \right\} \end{aligned} \quad (1)$$

where ρ is the continuous-time discount rate, $u(c) = c^{1-\gamma}/(1-\gamma)$ is CRRA utility with $\gamma = 1.5$, c_c^* is the child's equilibrium consumption policy, taken as given by the parent, and $\tilde{r}_c$ is the child's effective interest rate, which accounts for student debt (defined in [Section 2.3.5](#)). Consumption is bounded below by a floor $\underline{c} = \$3,900$ (in 2016 dollars), representing the effective consumption guaranteed by means-tested government transfers (SSI, SNAP, Medicaid), in the range estimated by [De Nardi et al. \(2010\)](#). The terms involving V_z^p and V_{zz}^p capture the effect of the child's income uncertainty on the parent's value function, while $V_{z_p}^p$ and $V_{z_p z_p}^p$ carry the parent's own persistent labor-income risk—the Ornstein–Uhlenbeck process for z_p of [Section 2.2](#), with mean-reversion $\kappa_p \equiv \kappa_{e_p}$ and diffusion $\sigma_p \equiv \sigma_{e_p}$, and where $r_{\text{eff}}(a_p)$ is the wealth-dependent return defined there. The final term is the lump-sum transfer channel: at Poisson arrivals with intensity λ_g , the parent may execute a one-time wealth rebalancing τ_L^* toward the child.

2.3.2 Child's Problem

The child chooses consumption c_c to maximize own lifetime utility:

$$\begin{aligned} \rho V^c = \max_{c_c} & \left\{ u(c_c) + V_{a_c}^c [\tilde{r}_c a_c + y_c - c_c + \tau^*] + V_{a_p}^c [r_{\text{eff}}(a_p) a_p + y_p - c_p^* - \tau^*] \right. \\ & + V_z^c(-\kappa z) + \frac{1}{2} \sigma^2 V_{zz}^c + V_{z_p}^c(-\kappa_p z_p) + \frac{1}{2} \sigma_p^2 V_{z_p z_p}^c + \frac{1}{2} \sigma_a^2 a_p^2 V_{a_p a_p}^c + \frac{1}{2} \sigma_a^2 a_c^2 V_{a_c a_c}^c + V_t^c \\ & \left. + \lambda_g [V^c(a_p - \tau_L^*, a_c + \tau_L^*, z, z_p; t) - V^c(a_p, a_c, z, z_p; t)] \right\} \end{aligned} \quad (2)$$

where τ^* and c_p^* are the parent's equilibrium policies, taken as given, and the jump term reflects that the child's state moves with the *parent's* lump-sum policy τ_L^* at transfer opportunities—altruism is one-sided, so the child receives but does not choose. The child tracks parent wealth a_p because it affects expected future transfers: wealthier parents are more likely to provide support.

2.3.3 Transfer Decision

The parent's optimal transfer is characterized by a complementarity condition:

$$\tau^* \geq 0, \quad V_{a_p}^p \geq V_{a_c}^p, \quad \tau^* (V_{a_p}^p - V_{a_c}^p) = 0 \quad (3)$$

This partitions the state space into two regions. In the *no-transfer region* ($\mathcal{N}$), the parent's marginal value of own wealth exceeds the marginal value of child wealth ($V_{a_p}^p > V_{a_c}^p$), and $\tau^* = 0$. In the *transfer region* ($\mathcal{T}$), these marginal values are equalized ($V_{a_p}^p = V_{a_c}^p$) and $\tau^* > 0$. The boundary between $\mathcal{N}$ and $\mathcal{T}$ —the *transfer boundary* $\partial\mathcal{T}$ —is determined endogenously as part of the equilibrium (Barczyk and Kredler, 2014a,b). The parent transfers when the child is relatively poor (low a_c), faces negative productivity shocks (low z), or when the parent is relatively wealthy (high a_p).

This complementary-slackness condition is the continuous-time analog of the transfer decision in discrete-time dynastic models, but with two key advantages. First, the flow transfer carries no within-period commitment: a discrete-time transfer is implicitly a promise held fixed for the length of the period, whereas the flow can be revised at every instant. Second, the boundary of the transfer region provides a clean characterization of the states

where moral hazard is operative. Flow transfers are not the only instrument: at stochastically arriving opportunities the parent can also move wealth in discrete amounts, and the same transfer boundary governs when it chooses to do so.

Transfer size. Within the transfer region, the transfer is determined at the consumption margin. The static optimality condition $u'(c_p) = \eta u'(c_c)$ implies an altruistic consumption target $c_c = \eta^{1/\gamma} c_p$ for the child, and the parent funds the gap between this target and the child's own flow resources $R_c = \tilde{r}_c a_c + y_c$, capped at the consumption level c_c^{opt} the child would choose for itself from its first-order condition:

$$\tau^* = \max\left\{0, \min\left(c_c^{\text{opt}} - R_c, \eta^{1/\gamma} c_p - R_c\right)\right\}. \quad (4)$$

The child consumes $R_c + \tau^*$ inside the transfer region: the transfer is consumption support, not a wealth transfer. The cap at c_c^{opt} reflects this—any gift beyond the child's desired consumption would be saved by the child, converting parent wealth into child wealth.⁷

The flow transfer never moves the parent's wealth stock. Mechanically, moving a discrete amount in an instant would require an unbounded flow; economically, the absence of commitment makes pre-funding unattractive, since wealth placed in the child's account would be consumed too quickly. By retaining the wealth the parent preserves the option to support the child only when it is actually constrained, providing liquidity period by period rather than capital while continuing to optimize its own saving inside the transfer region ($\dot{a}_p = r_{\text{eff}}(a_p)a_p + y_p - c_p - \tau^*$). Appendix F derives these properties formally.

2.3.4 Lump-Sum Transfer Opportunities

The flow transfer is consumption support by construction—capped at the child's desired consumption, it props up the child's spending in bad states but never adds a dollar to the child's balance sheet. Observed inter-vivos giving is not so confined: parents make discrete, sizeable transfers—help with a house down payment, a wedding, support around

⁷Transfers are therefore not confined to children exactly at the borrowing constraint, as in the no-labor-income setting of [Barczyk and Kredler \(2014b\)](#) where gifts flow only to a recipient with zero wealth. With flow income, the gift (4) is positive whenever the child's flow resources fall short of both its desired consumption and the altruistic target, $R_c < \min(c_c^{\text{opt}}, \eta^{1/\gamma} c_p)$, so the transfer region concentrates at low child wealth and adverse income draws but its boundary lies at strictly positive a_c (Figure 2).

the birth of a grandchild—that move *wealth* between generations well before any bequest. In the HRS, transfers of \$500 or more in a two-year window occur for 5 to 25 percent of parent-child pairs depending on parent wealth, and their lumpiness is intrinsic to how they arrive: they are occasioned by life events, not paid as annuities. A model whose only wealth-moving instruments are a tuition gift at age 18 and a bequest at the parent’s death cannot reproduce the lumpy inter-vivos transfers observed between enrollment and death—the very transfer incidence the extensive-margin moments measure—with an instrument set that, by construction, produces none of them in between.

I restore the lump without restoring the commitment.⁸ At exogenous Poisson arrivals with intensity λ_g , the parent may execute a one-time transfer $\tau_L \in [0, \max(a_p, 0)]$, chosen at the moment of the opportunity to maximize its own value,

$$\tau_L^*(a_p, a_c, z, z_p; t) = \arg \max_{\tau_L \in [0, \max(a_p, 0)]} V^p(a_p - \tau_L, a_c + \tau_L, z, z_p; t). \quad (5)$$

Because the arrival is exogenous, the parent chooses neither the timing nor a schedule: the exercise is an instantaneous action, not a promise, so no commitment is violated; and because opportunities arrive smoothly in time, the channel introduces none of the anticipation discontinuities that deterministic transfer dates would create (Barczyk and Kredler, 2021). Proposition 1 in Appendix H makes this precise: under the Poisson channel the value functions remain C^2 away from the transfer boundary, whereas deterministic dates would reintroduce the non-smoothness. The arrival intensity has a natural interpretation as the frequency of life events at which discrete transfers occur.

The optimality condition for (5) ties the lump to the same transfer boundary that governs the flow. An interior $\tau_L^* > 0$ requires $V_{a_c}^p > V_{a_p}^p$ at the pre-transfer state—strictly inside the transfer region—and the optimal lump moves the state onto the boundary where the two marginal values are equalized. Lump-sum transfers are therefore infrequent, sizeable, and state-contingent: they fire when the child has fallen deep into the transfer region, and they rebalance the dynasty’s wealth in one shot rather than drip-feeding consumption. This is

⁸There is also a methodological reason this channel is absent from continuous-time altruism models. Barczyk and Kredler (2014a) observe that discrete-time dynastic models implicitly assume the donor can commit to a lump-sum transfer for the length of a period—commitment smuggled in by the time aggregation itself—and the continuous-time reformulation removes it; but in removing the artificial commitment it also removes the lump, leaving only the flow.

precisely the empirical object the extensive-margin transfer moments measure—and it is what identifies the altruism parameter η , since the wealth gradient in transfer incidence pins down the size of the region in which the parent exercises the opportunity—giving altruism a channel through which it can *raise* child wealth in adverse states rather than only eroding the child’s own saving through the Samaritan’s dilemma. Routing general altruistic wealth rebalancing through these opportunities also keeps it off the college gift: the mean-support moment disciplines the college gift at enrollment, while the transfer-probability moments by parent wealth discipline the Poisson lumps over the rest of the overlap, so the two margins are separately identified rather than loaded onto a single college transfer. The remaining lump-sum transfers in the model are the college-conditional gift at enrollment and the bequest at death; lump opportunities begin at the child’s college-exit age, so the college gift remains the only lump instrument during the schooling years.

2.3.5 Consumption Policies and Wealth Dynamics

The first-order conditions for consumption are:

$$u'(c_p^*) = V_{a_p}^p, \quad (6)$$

$$u'(c_c^*) = V_{a_c}^c. \quad (7)$$

Under CRRA utility, these invert to $c_p^* = (V_{a_p}^p)^{-1/\gamma}$ and $c_c^* = (V_{a_c}^c)^{-1/\gamma}$. The drifts of the two wealth processes (the deterministic part of the laws of motion $da_p = \dot{a}_p dt + \sigma_a a_p dB_p$ and $da_c = \dot{a}_c dt + \sigma_a a_c dB_c$) are:

$$\dot{a}_p = r_{\text{eff}}(a_p) a_p + y_p(t) - c_p - \tau, \quad (8)$$

$$\dot{a}_c = \tilde{r}_c a_c + y_c(t, z) - c_c + \tau, \quad (9)$$

where the parent’s wealth is bounded below by the age- and education-dependent borrowing limit introduced above (and zero from retirement on). The child’s wealth, by contrast, can be negative: a child who attends college may finance tuition and consumption with student loans, accumulating a debt balance ($a_c < 0$) that it carries into working life (Section 2.7).

The child's effective interest rate therefore depends on the sign and source of its debt,

$$\tilde{r}_c = r + \iota_s \mathbf{1}\{\text{student debt}\} + \iota_b \mathbf{1}\{\text{unsecured debt}\},$$

so the child earns r on savings, pays the fixed student-loan rate $r + \iota_s$ on student debt, and pays the estimated wedge $r + \iota_b$ on other unsecured debt. During college the constraint is relaxed to $a_c \geq -\bar{b}$ at the student rate, and that balance is amortized so that $a_c \geq 0$ by age 32 (the standard ten-year repayment plan). As a working-age adult the child—like the parent—may also hold unsecured debt up to the education-specific limit (\$25,000 for high school, \$85,000 for college) at the rate $r + \iota_b$, with the limit ramping to zero before retirement; this borrowing generates the negative net worth observed among young adults.

2.4 Child-Alone Problem and the Dynastic Link

When the parent dies (parent age 72, child age 48), the child solves a consumption-savings problem with their own accumulated wealth—augmented by the bequest—for the remainder of their life:

$$\begin{aligned} \rho V^{c,\text{alone}} = \max_{c_c} & \left\{ u(c_c) + V_{a_c}^{c,\text{alone}} [\tilde{r}_c a_c + y_c(t, z) - c_c] \right. \\ & \left. + V_z^{c,\text{alone}}(-\kappa z) + \frac{1}{2} \sigma^2 V_{zz}^{c,\text{alone}} + \frac{1}{2} \sigma_a^2 a_c^2 V_{a_c a_c}^{c,\text{alone}} + V_t^{c,\text{alone}} \right\} \end{aligned} \quad (10)$$

The child lives alone until its own death at age 72, working until 66 and then drawing social security. The dynasty is then closed as a stationary recursion: a new dynasty begins with a parent whose initial wealth equals the child's terminal wealth a_c at age 72, whose education is $e'_p = e_c$, and whose child has ability θ' drawn from the intergenerational transmission process.

2.5 Terminal Conditions

At the parent's death ($t = T$, parent age 72, child age 48), the coupled system terminates and the child transitions to the child-alone problem. The boundary conditions at $t = T$ are:

$$V^p(a_p, a_c, z, z_p; T) = \phi_b \frac{(a_p + \underline{a}_{\text{beq}})^{1-\gamma}}{1-\gamma} + \eta \cdot V^{c, \text{alone}}(a_c + \alpha_{\text{beq}} a_p, z; 0), \quad (11)$$

$$V^c(a_p, a_c, z, z_p; T) = V^{c, \text{alone}}(a_c + \alpha_{\text{beq}} a_p, z; 0). \quad (12)$$

The child's value function transitions to the child-alone problem (10), evaluated at $t = 0$ of the child-alone clock—the moment of the parent's death. Equation (12) states this continuity: the child carries forward its own wealth a_c plus the bequest $\alpha_{\text{beq}} a_p$.

The parent's terminal value (11) has two components. The first, $\phi_b (a_p + \underline{a}_{\text{beq}})^{1-\gamma}/(1-\gamma)$ with $\phi_b \geq 0$, is a warm-glow term following De Nardi (2004) that captures the parent's direct utility from holding wealth at the end of the overlap period.⁹ The second component, $\eta \cdot V^{c, \text{alone}}(a_c + \alpha_{\text{beq}} a_p, z; 0)$, reflects the parent's altruistic concern for the child's continuation utility, with $\alpha_{\text{beq}} \in [0, 1]$ the fraction of parent terminal wealth bequeathed at death. The bequest channel creates two motives for wealth accumulation—warm glow (ϕ_b) and altruism through α_{beq} , which disciplines the intergenerational wealth correlation.

2.6 College Decision

At the beginning of the dynasty ($t = 0$), two decisions are taken in sequence: the parent makes a one-time inter-vivos gift, and the child then decides whether to attend college. Let $V_e^c(a_p, a_c)$ and $V_e^p(a_p, a_c)$ denote the child's and parent's equilibrium continuation values under education $e \in \{C, HS\}$, evaluated at $z = 0$.

Child's decision. Given the parent's college gift (described below), the child weighs its own value gain from attending—the college branch evaluated at the post-gift state, the high-

⁹The shifter $\underline{a}_{\text{beq}} \geq 0$ makes bequests a luxury good in the sense of De Nardi (2004): with $\underline{a}_{\text{beq}} > 0$ the marginal propensity to bequeath rises with wealth, so poor parents leave little while the rich bequeath a rising share—the mechanism that generates the fat right tail of the wealth distribution.

school branch at the no-gift state—

$$G \equiv V_C^c - V_{HS}^c,$$

against the psychic cost of college, which follows the [Abbott et al. \(2019\)](#) linear-in-log specification with an idiosyncratic shock:

$$\kappa(\theta, \varepsilon_\kappa, e_p) = (\kappa_0 + \kappa_\theta \log \theta + \kappa_w \mathbf{1}\{e_p = C\} + \sigma_\kappa \varepsilon_\kappa) \cdot \bar{V}, \quad \varepsilon_\kappa \sim N(0, 1),$$

where ε_κ is drawn once per individual at $t = 0$ and $\bar{V} = \text{mean}_{a_p, \theta} |V_C^c - V_{HS}^c|$ is a normalization constant that makes the dimensionless parameters κ_0 , κ_θ , κ_w , and σ_κ comparable in scale to the college value gap. The gradient $\kappa_\theta < 0$ implies that higher-ability individuals face lower psychic costs, reflecting the non-pecuniary, “consumption value” view of schooling ([Cunha et al., 2005](#); [Heckman et al., 2006](#)). The shifter $\kappa_w < 0$ lowers the psychic cost for children of college-educated parents, a non-pecuniary intergenerational channel of college taste, information, and preparation that operates through the human-capital preparedness parental education builds over childhood ([Cunha and Heckman, 2007](#); [Cunha et al., 2010](#)).¹⁰ It lets the socioeconomic gradient in attainment operate through the *non-pecuniary* cost of college, distinct from the credit-market frictions emphasized by [Belley and Lochner \(2007\)](#) and [Caucutt and Lochner \(2020\)](#), which the model captures separately through the borrowing channels of Section 2.7. This is the only channel through which parent education enters the college decision beyond altruism and resources; the parent derives no separate utility from the child’s enrollment. The child attends whenever the value gain exceeds the realized cost, $G \geq \kappa(\theta, \varepsilon_\kappa, e_p)$. Integrating over the normally distributed cost shock yields a probit college-

¹⁰The socioeconomic gradient in attainment could reflect this child-side taste/preparation channel (κ_w) or a parent-side paternalistic motive (ξ , as in [Abbott et al. 2019](#)), in which the parent values the child’s enrollment independently of the child’s own preference. The two are not separately identified by the attendance gradient alone, but they differ in their transfer implications: paternalism operates only through the parent’s college gift, so a paternalism-driven gradient implies counterfactually large inter-vivos transfers, which the transfer moments do not show. The baseline therefore omits the paternalistic motive and attributes the gradient to the child-side shifter κ_w .

attendance probability:

$$\Pr(e_c = C \mid a_p, a_c, \theta, e_p) = \begin{cases} \Phi\left(\frac{G/\bar{V} - \kappa_0 - \kappa_\theta \log \theta - \kappa_w \mathbf{1}\{e_p = C\}}{\sigma_\kappa}\right) & \text{if } G > 0, \\ 0 & \text{otherwise,} \end{cases} \quad (13)$$

where Φ is the standard normal CDF and σ_κ governs the dispersion of college choices: a larger σ_κ flattens the mapping from the value gap to attendance and, through favorable cost draws, raises enrollment among low-ability types. The decision is the child's alone—consistent with one-sided altruism, the child does not weigh the parent's value in choosing its education.

Parent's college gift. Before the child's cost shock is realized, the parent chooses a one-time college-conditional gift τ_C , paid at enrollment, anticipating how it shifts the child's choice:

$$\begin{aligned} \tau_C^* = \arg \max_{\tau_C \in [0, a_p]} & \Pr(C \mid \tau_C) V_C^p(a_p - \tau_C, \tau_C) \\ & + (1 - \Pr(C \mid \tau_C)) V_{HS}^p(a_p, 0), \end{aligned} \quad (14)$$

where $\Pr(C \mid \tau_C)$ is the probit (13) with the college branch evaluated at the post-gift state $(a_p - \tau_C, a_c = \tau_C)$ and the high-school branch at $(a_p, 0)$. The conditional gift requires no commitment: payment and enrollment are a single transaction—the parent finances enrollment at the moment it occurs—so there is no promise to renege on, and a child who chooses high school receives no lump sum at $t = 0$. The gift is bounded only by the parent's wealth, $\tau_C \in [0, a_p]$: it is the conditional inter-vivos transfer the parent makes at enrollment, the broad earmarked support—tuition together with allowances and in-kind room and board—that the college-support data of [Abbott et al. \(2019\)](#) and the mean-support moment of Section 4 measure, rather than tuition narrowly. The size of the gift is therefore disciplined directly by that mean-support moment. The gift is purely altruistic: it raises the child's continuation value V_C^p that the parent already internalizes, and the parent derives no additional utility from enrollment per se.¹¹ Wealthy parents—whose marginal utility of consumption is low—make larger gifts, relaxing the child's college-financing constraint and raising the prob-

¹¹An enrolling child enters the overlap with $a_c = \tau_C^*$ and the parent with $a_p - \tau_C^*$; a child who chooses high school receives no lump sum at $t = 0$.

ability of enrollment; together with the parent-education cost shifter κ_w , this is the force that generates college attendance among children of high-resource and college-educated families. During the schooling years the college gift is the only lump-sum instrument; after college exit, support moves through the state-contingent flow $\tau(t)$, the Poisson lump-sum opportunities of Section 2.3.4, and ultimately the bequest.

Parental altruism enters the college decision through the continuation values V_C and V_{HS} . Parents who anticipate supporting their child during and after college raise the value of attending, and because the continuation values incorporate the entire future trajectory of transfers—which depends on the child’s income path and hence on education—parent wealth affects enrollment beyond the direct cost channel. Anticipated transfers, however, raise both V_C and V_{HS} : they subsidize college enrollment but also insure non-college children against income risk. The net effect of the Samaritan’s dilemma on the college decision is therefore theoretically ambiguous; Section 6.1 formalizes this result.

2.7 College Financial System

I model the college financial system following [Abbott et al. \(2019\)](#); Appendix G gives the full specification. The annual net cost of college is $\phi = \$6,700$ (average net tuition from the NLSY97), and students work a fraction $\ell_C = 0.56$ of full-time hours, earning at the high-school rate until graduation. Financial aid follows a continuous expected-family-contribution schedule based on parental *income* at college entry: the maximum grant $\bar{g}$ phases out linearly in parent income, so the effective price of college rises with family resources. During college the child may borrow up to a cumulative limit

$$\bar{b}(a_p) = \bar{b}_0 + b_{\text{slope}} \cdot a_p$$

at the student-loan rate $r + \iota_s$, where the base limit $\bar{b}_0 = \$18,000$ is the federally available capacity and the collateral channel $b_{\text{slope}} \geq 0$ adds private/PLUS capacity requiring parental backing, so children of wealthier parents can borrow more; the balance amortizes to $a_c \geq 0$ by age 32. As [Abbott et al. \(2019\)](#) document, parental transfers supplement institutional aid and government grants can crowd out private contributions.

2.8 Equilibrium

A Markov-Perfect Equilibrium (MPE) consists of value functions V^p , V^c , consumption policies c_p^* , c_c^* , a transfer-flow policy τ^* , a lump-sum transfer policy τ_L^* , and a college enrollment rule e_c^* , for each (θ, e_p, e_c) pair, such that: (i) given c_c^* , τ^* , and τ_L^* , the parent’s HJB (1), the complementary-slackness condition (3), and the lump-sum optimality condition (5) are satisfied; (ii) given c_p^* , τ^* , and τ_L^* , the child’s HJB (2) is satisfied; (iii) after the parent’s death ends the overlap ($t > T$), $V^{c,\text{alone}}$ satisfies (10); (iv) the college rule follows from (13); and (v) boundary conditions (11)–(12) and borrowing constraints hold at all times.

The equilibrium is solved backward: first the child-alone problem, then the coupled system during the overlap period, and finally the college choice at $t = 0$. Within the coupled system, the transfer is obtained in closed form at each state—the bounded gift (4)—so the transfer region is determined pointwise rather than by iterating on the complementary-slackness condition, and the value functions are updated by implicit time-stepping. The model is dynastic in that generations are *linked* rather than fully recursive in the value functions: altruism is embedded one generation forward—the parent values the child through $\eta V^{c,\text{alone}}$, while the child’s post-overlap life is solved as a standalone life-cycle problem that does not itself carry the child’s altruism toward a future child—and successive generations are connected through the simulation, where each child’s terminal wealth, ability, and education initialize the next generation’s parent and generate a stationary cross-generation distribution. I therefore solve the value functions once, by backward induction over a single parent–child cycle, rather than iterating them across generations to a recursive dynastic fixed point as in Barczyk et al. (2023); this is a deliberate simplification: the Samaritan’s-dilemma mechanism operates within the parent–child relationship, and the child’s own altruism toward the next generation is second order for the college decision at $t = 0$, which turns on the child’s lifetime resources and the support it anticipates from its parent. Appendix F derives the equilibrium properties, including the Samaritan’s dilemma distortion in continuous time, and Appendix H describes the numerical solution algorithm.

2.9 Empirical Implications

College reshapes the transfer region—the set of states in which the parent optimally supports the child—and the boundary shift maps into family outcomes observable in matched parent-child data.

Why college shrinks the transfer region. Recall that the parent transfers when the marginal value of own wealth equals the marginal value of child wealth ($V_{a_p}^p = V_{a_c}^p$). Because the parent weighs the child’s utility by η in the objective (1), $V_{a_c}^p$ is increasing in η : more altruistic parents have a larger transfer region. This condition is more likely to bind when the child is relatively poor or faces negative income shocks. College raises the child’s expected income through a higher ability gradient ($\lambda_C > \lambda_{HS}$) and a steeper life-cycle earnings profile; the resulting increase in child wealth pushes states away from the transfer boundary, so college-educated children spend less time receiving support and the region contracts relative to high-school children. Figure 2 illustrates this in the (a_c, a_p) state space: college shifts the boundary to the left, so that at any given level of parent wealth the child must be poorer before transfers become optimal. The gap between the two boundaries—the states in which a high-school child would receive transfers but a college child would not—is the contraction of the region; together with the overconsumption zone bordering it, where the child under-saves in anticipation of entering the region, this contraction captures the scope for moral hazard that college eliminates.




Figure 2. Transfer Region Boundary: College vs. High School

Notes: Schematic illustration of the transfer region in the (a_c, a_p) state space. In the transfer region (shaded), the parent optimally transfers to the child ($\tau > 0$). The solid boundary corresponds to a college-educated child; the dashed boundary to a high-school-educated child. College shrinks the region because higher expected income keeps child wealth further from the transfer boundary. The model-computed transfer boundary that this schematic illustrates is reported in Figure 4, and the full transfer policy in Appendix F (Figure F.1).

The contraction of the transfer region yields two predictions, both conditional on parent income, that I document in the descriptive evidence of Section 3. The first is a consumption release: a smaller region means parents spend less of their lifetime making transfers, so resources that would have flowed to the child under the high-school path are instead available for own consumption, and parent consumption is higher when the child attends college. The second concerns the two margins of transfers. Because a larger negative shock is required to push a college child into the region, the probability of transferring falls; conditional on transferring, the states reached lie further inside it, so the optimal transfer τ^* is weakly

larger. The model thus predicts a negative extensive-margin effect and a zero-or-slightly-positive intensive-margin effect.

Parents hold wealth against the option of supporting a child who falls into the transfer region, so college releases resources even in states where no transfer is ever paid. The consumption release therefore reflects both fewer realized transfers and less precautionary saving, and the model predicts that the first exceeds the second: the consumption gain from college is larger than the fall in realized transfers. This is the continuous-time option-value channel of [Barczyk and Kredler \(2014b\)](#), in which the family inefficiency operates before transfers flow rather than through the transfers themselves. The estimated model quantifies this channel directly.

3 Empirical Evidence

This section documents how parent consumption and inter-vivos transfers vary with the child’s college attainment in matched parent-child data (PSID, HRS, and NLSY97). Each fact compares outcomes for parents whose children attended college with those whose children did not, conditioning on parent income to absorb the parent’s own resources. None of these moments is targeted in estimation, so [Section 5.2](#) confronts the estimated model with each of them as an out-of-sample check.

3.1 Data

The empirical analysis draws on three data sources. The Panel Study of Income Dynamics (PSID) provides matched parent-child panels with consumption, income, wealth, and inter-vivos transfer data from 1999 onward. I link parents to adult children using the Family Identification Mapping System (FIMS). The summary-statistics sample ([Appendix Tables B.1–B.2](#)) comprises 53,326 parent-year observations from 8,362 households with heads aged 25–60; the parent-consumption analysis of [Section 3.2](#) further restricts to parents aged 50+ with adult children aged 26+, collapsed to the parent-year level, yielding 15,830 observations. The Health and Retirement Study (HRS, 1992–2014 Family-Kids file) offers a larger sample of older parent-child pairs with detailed transfer information—548,239 parent-to-child wave observations, of which 78,541 involve a positive transfer—which I use to study the ex-

tensive and intensive margins of inter-vivos transfers. The National Longitudinal Survey of Youth 1997 (NLSY97) provides cognitive ability scores (ASVAB) and parental wealth for 5,400 individuals.

3.2 Parent Consumption and Child College Status

The model implies that, conditional on own income and wealth, parents of college children consume more (the consumption release of Section 2.9): college shifts children out of the transfer region by permanently raising their income, freeing parents from contingent obligations even when no transfers flow.

I examine this implication by estimating the following regression on the PSID sample at the parent-year level:

$$C_{p,t} = \beta_0 + \beta_1 \text{ShareCollege}_{p,t} + \beta_2 N_{p,t} + \delta_{Q_p^{inc}} + \beta_X \mathbf{X}_{\mathbf{p},t} + \varepsilon_t + \epsilon_{p,t}$$

where $C_{p,t}$ is household consumption (in 2016 dollars) of parent p at time t , $\text{ShareCollege}_{p,t}$ is the fraction of parent p 's children with a bachelor's degree or higher, $N_{p,t}$ is the total number of children, $\delta_{Q_p^{inc}}$ are parent income quartile fixed effects (within-cohort ranking), and $\mathbf{X}_{\mathbf{p},t}$ is a comprehensive set of controls: parent total wealth, non-financial wealth, household income, wealth quartile, labor force participation, household size, head's birth year, education, U.S. state, a cubic in head's age, housing tenure, and race. The specification includes year fixed effects ε_t .

A higher share of college-educated children is associated with higher parent consumption: moving from none to all of a parent's children holding a BA raises consumption by about 8% (Table 1, Column 1). The association is robust to alternative definitions of college exposure—having any child with a BA is associated with 8% higher consumption (Column 2), and having all children hold a BA with 6% higher consumption (Column 3), both significant at the 1% level. In levels (Column 4), moving from no college children to all corresponds to roughly \$2,500 more in annual parent consumption.

Table 1. Parent Consumption and Child College Status (PSID)

	(1)	(2)	(3)	(4)
	Log C	Log C	Log C	Level C (000s)
Share of Children with BA+	0.081*** (0.019)			2.470*** (0.866)
Number of Children	0.014* (0.008)	0.010 (0.008)	0.016* (0.008)	0.170 (0.287)
Any Child with BA+		0.080*** (0.017)		
All Children with BA+			0.056*** (0.017)	
Controls	Yes	Yes	Yes	Yes
Year FE	Yes	Yes	Yes	Yes
Parent income Q FE	Yes	Yes	Yes	Yes
Observations	15,830	15,830	15,830	15,831
R^2	0.558	0.558	0.557	0.357

Notes: OLS regressions of parent household consumption on child college indicators. “Share of Children with BA+” is the fraction of children with a bachelor’s degree or higher. The dependent variable in Column 4 is total consumption in thousands of 2016 dollars. Controls include a cubic in parent age, household income, non-housing wealth, total wealth, wealth quartile, and fixed effects for year, education, parent income quartile, gender, housing tenure, state, labor force status, race, household size, and birth year. Standard errors clustered at the parent household level in parentheses. * $p < .10$, ** $p < .05$, *** $p < .01$.

3.3 Inter-Vivos Transfers

If college frees parents from contingent obligations, do realized transfers fall correspondingly? Answering requires decomposing transfers into their extensive and intensive margins, which map directly onto the model’s transfer-region structure (Section 2.9).

Most parent–child pairs report zero transfers in any given wave: positive transfers occur in only 14.4% of pair-wave observations overall—12.5% for non-college and 18.7% for college children (Table 2). Conditional on a transfer, amounts are substantial and right-skewed: the median is \$2,413 for non-college children and \$3,672 for college children, with 90th percentiles

of \$14,135 and \$24,481 respectively. This mass at zero with a heavy right tail is what the model predicts: the transfer region is entered infrequently, but entry is a non-trivial insurance event. In these raw cross-sectional comparisons college children receive transfers *more* often and in larger amounts—unsurprisingly, since college children come from wealthier families—so recovering the within-family effect of college requires the parent fixed effects I turn to next.

Table 2. Distribution of Inter-Vivos Transfers from Parents to Children, by Child College Status (HRS)

Pr(Transfer > 0)	0.144
E(Transfer Transfer > 0)	\$8,161
Median(Transfer Transfer > 0)	\$2,897
P75(Transfer Transfer > 0)	\$7,968
P90(Transfer Transfer > 0)	\$18,394
E(Transfer, unconditional)	\$1,172
Observations	548,239
Observations with Transfer > 0	78,541

Notes: Wave-level statistics for parent-to-child transfers from HRS 1992–2014 (Family-Kids file, waves 1–12). Sample restricted to parent age ≥ 50 , child age ≥ 26 . Transfer amounts in real 2016 US dollars. Unconditional mean includes zeros.

I decompose the within-family college effect along these two margins. Appendix Table C.1 reports linear probability models for $\text{Pr}(\text{Transfer} > 0)$ with parent fixed effects:

$$\mathbf{1}[IVT_{ijt} > 0] = \beta_0 + \beta_1 \text{College}_j + \gamma_t + \beta_X \mathbf{X}_{\mathbf{j}t} + \delta_{Q_i^p} + \varepsilon_i + \epsilon_{ijt}$$

College is associated with a 3.1 percentage-point lower probability of any transfer—a 21% reduction on the 14.4% mean (Appendix Table C.1, Column 1). Interacting college with parent income quartile (Column 2) shows a steep gradient: -1.0 pp for Q1, -2.8 for Q2, -3.4 for Q3, and -4.2 for Q4. Wealthier parents have larger transfer regions *ex ante*, so college contracts them by more.

Conditional on a positive transfer, college children receive \$557 more per event (Appendix Table C.2, Column 1, $p < 0.10$). The effect is heterogeneous (Column 2): $-\$906$ for Q1 parents but positive and large above— $\$958$ for Q2 and $\$826$ for Q4. This matches the boundary shift: college moves the transfer boundary left in (a_c, a_p) space, so the states that pull a college child into the region involve worse income draws; conditional on entry the child is further inside the region and the optimal transfer is larger, most so for wealthier parents whose boundary shift is more pronounced.

The 3.1-point fall in transfer probability corresponds to roughly \$253 per year in expected transfers ($0.031 \times \$8,161$), an order of magnitude below the \$2,500 consumption difference of Section 3.2. The comparison is suggestive, but it is the option-value pattern: college lowers the burden of standing ready to support the child, releasing resources beyond realized transfers. Altogether, college is associated with higher parent consumption and less frequent but conditionally larger transfers.

4 Estimation

I estimate the model in two stages. In the first stage, I set parameters that are well-identified from external data or the existing literature, including the income-process parameters, which I take from Abbott et al. (2019) and map from their discrete-time estimates to the continuous-time counterparts used in the model. In the second stage, I estimate the remaining thirteen structural parameters by the simulated method of moments (SMM), matching 30 targeted moments.

4.1 Externally Set Parameters

Table 3 reports the parameters set in the first stage. The coefficient of relative risk aversion is $\gamma = 1.5$, following Abbott et al. (2019). The annual real interest rate is $r = 0.03$, following Daruich and Kozlowski (2020). The continuous-time discount rate ρ is estimated internally by SMM to match the wealth-to-income ratio and wealth distribution moments. The college costs, borrowing limits, and financial-aid schedule parameters $(\bar{g}, \bar{y}_{\text{efc}})$ are set externally to the Abbott et al. (2019) values, as described in Section 2.7. The income-process parameters—the education-specific ability gradients λ_e , age profiles, OU mean-reversion rates κ_e and

diffusions σ_e , and entry dispersions—and the intergenerational ability parameters $(\rho_\theta, \sigma_\theta, \mu_\theta)$ are likewise taken from [Abbott et al. \(2019\)](#). The relative price of college labor ω_C is calibrated externally to target the high-school-to-college income ratio (Section 4), and the student-loan rate wedge ι_s is fixed; the wedge on working-age unsecured borrowing ι_b is likewise fixed externally at the [Abbott et al. \(2019\)](#) value ($\iota_b = 0.064$). The multiplicative wealth-diffusion volatility is fixed at $\sigma_a = 0.10$ per year, following [Barczyk and Kredler \(2014a\)](#). The lump-sum transfer arrival intensity is set externally at $\lambda_g = 0.25$ per year, so a parent has on average one opportunity to make a lump-sum transfer every four years. These opportunities arrive at random, independently of the family’s wealth, so λ_g governs only how often a transfer *can* occur, not how transfer behavior varies with wealth. The wealth gradient in transfers instead comes from whether the parent chooses to exercise an opportunity—the size of the transfer region, which altruism η governs. Fixing λ_g therefore keeps the identification of η from the transfer-probability moments clean. Retirement income is estimated from HRS data as the person-weighted average of annual Social Security and employer-pension income among respondents aged 66–72 receiving benefits, computed separately for college (bachelor’s degree or higher) and non-college parents. Finally, the intergenerational income-transmission parameter ρ_{zp} , which governs the persistence of the parent’s residual income type into the child’s initial productivity net of the transmitted-ability and parental-college channels, is not separately identified from those channels by the targeted moments; I therefore fix it at 0.23, a conservative value for the component of intergenerational earnings persistence—a rank–rank slope of 0.34 in [Chetty et al. \(2014\)](#) and an intergenerational elasticity of roughly 0.4 in [Solon \(1992\)](#)—not already carried by transmitted ability (ρ_θ) and the parental-college earnings premium (ζ_{pe}). The bequest fraction is fixed externally on the same logic at $\alpha_{\text{beq}} = 0.34$, the bequeathed share of the estate, with the complement absorbed by end-of-life medical and long-term-care expenses ([De Nardi et al., 2010](#); [Lockwood, 2018](#)). A robustness check confirms that perturbing α_{beq} and ρ_{zp} leaves the targeted moments and the counterfactuals essentially unchanged.

Table 3. Externally Set Parameters

Parameter	Description	Value	Source
<i>Preferences</i>			
r	Real interest rate	0.03	Daruich and Kozlowski (2020)
γ	CRRA risk aversion	1.5	Abbott et al. (2019)
$\underline{c}$	Consumption floor (\$000s/yr)	3.9	De Nardi (2004)
σ_a	Wealth-return volatility (/yr)	0.10	Barczyk and Kredler (2014a)
a_{ref}	Return-premium saturation scale (\$000s)	100	Normalization
λ_g	Lump-transfer arrival rate (/yr)	0.25	Normalization
<i>Demographics (years)</i>			
T^{ret}	Parent working, overlap (child 18–42)	24	—
$T - T^{\text{ret}}$	Parent retired, overlap (child 42–48)	6	—
T_{alone}	Child alone, working (48–66)	18	—
$T_{\text{retire},c}$	Child retirement (66–72)	6	—
<i>Ability gradients</i>			
λ_C	Returns to ability, college	0.797	Abbott et al. (2019)
λ_{HS}	Returns to ability, HS	0.517	Abbott et al. (2019)
<i>Earnings process (annual)</i>			
ρ_C^{ann}	Persistence, college	0.966	Abbott et al. (2019)
ρ_{HS}^{ann}	Persistence, HS	0.952	Abbott et al. (2019)
$(\sigma_\eta^{\text{ann}})^2$	Innovation variance (both groups)	0.017	Abbott et al. (2019)
σ_C	OU diffusion, college	0.133	Abbott et al. (2019)
σ_{HS}	OU diffusion, HS	0.134	Abbott et al. (2019)
$\bar{w}$	Avg. annual earnings	\$70k	Census
<i>College costs</i>			
ϕ_C	Annual tuition (before aid)	\$6.7k	Abbott et al. (2019)
ℓ_C	Time worked in college	0.56	Census
<i>Intergenerational ability</i>			
ρ_θ	Persistence of log ability	0.36	Abbott et al. (2019)
σ_θ	Std. dev. of ability innovation	0.40	Abbott et al. (2019)
μ_θ	Mean of log ability	1.30	Abbott et al. (2019)
<i>Bequests & intergenerational transmission (fixed)</i>			
α_{beq}	Bequest fraction (share of estate to children)	0.34	De Nardi et al. (2010) ; Lockwood (2018)
ρ_{zp}	Intergen. income transmission ($z_p \rightarrow z_0$)	0.23	Calibrated
<i>Financial aid & borrowing</i>			
$\bar{g}$	Max. annual grant (\$000s)	2.82	Abbott et al. (2019)
$\bar{y}_{\text{efc}}$	EFC income phase-out (\$000s)	72.0	Abbott et al. (2019)
$\bar{b}_0$	Student loan limit (base)	\$18.0k	Federal
ι_s	Student-loan rate wedge (fixed)	0.031	Abbott et al. (2019)
ι_b	Unsecured borrowing rate wedge (fixed)	0.064	Abbott et al. (2019)
$\underline{a}_C$	Unsecured borrowing limit, college	\$85k	Abbott et al. (2019)
$\underline{a}_{HS}$	Unsecured borrowing limit, HS	\$25k	Abbott et al. (2019)
<i>Retirement income (annual)</i>			
SS_C	Soc. Security + pension, college (BA+)	\$27.9k	HRS
SS_{HS}	Soc. Security + pension, non-college	\$20.3k	HRS

Notes: All dollar amounts in 2016 dollars. The continuous-time income process is set externally from [Abbott et al. \(2019\)](#) (male estimates): persistence maps to the mean-reversion rate $\kappa_e = -\ln \rho_e^{\text{ann}}$, and the OU diffusion σ_e is chosen so the stationary variance matches AGMV (Appendix D). The transitory-income standard deviations $\sigma_{\varepsilon,e}$, the loan-limit slope b_{slope} , and the return-on-wealth gradient $r_{a,\text{grad}}$ are instead estimated by SMM (Table 5); σ_a regularizes the transfer boundary (Section 2). The financial-aid schedule $g(y_p^0) = \bar{g} \max(0, 1 - y_p^0/\bar{y}_{\text{efc}})$ keys on parent income at entry; it, the loan limit $\bar{b}_0$, and the rate wedges ι_s, ι_b follow [Abbott et al. \(2019\)](#). Retirement income is the person-weighted mean of annual Social Security plus private/employer pension among HRS respondents aged 66–72 receiving positive benefits, by education (college = bachelor’s degree or higher), in 2016 dollars. The lump-transfer arrival rate λ_g is a normalization (Section 2.3.4): the transfer-region geometry, not the arrival rate, governs the wealth gradient in transfer incidence.

4.2 Income Process Estimation

The income process in continuous time is an Ornstein-Uhlenbeck process $dz = -\kappa_e z dt + \sigma_e dW$ with education-specific parameters, set externally from the male estimates of [Abbott et al. \(2019\)](#), hereafter AGMV.¹² I map their discrete-time estimates to the continuous-time process and hold it fixed rather than estimating the diffusion coefficients within the structural routine.

In discrete time, log labor income decomposes as $\log \epsilon_j = \lambda_e \log \theta + \gamma_{e,j} + z_j$, with education-specific ability gradient λ_e , deterministic age profile $\gamma_{e,j}$, and persistent shock z_j . AGMV estimate the gradient from the NLSY79 (log hourly wages on log AFQT, by education group), the age profile from a PSID age quadratic, and the shock persistence by minimum distance on the autocovariance structure of wage residuals:

$$\begin{aligned} z_{iat}^e &= \log y_{it} - \widehat{f^e(a_{it})} - \hat{\lambda}_e \log(\text{AFQT}_i) \\ z_{iat}^e &= \rho_e^{\text{ann}} z_{i,a-1,t-1}^e + \eta_{iat}^e \\ \eta_{iat}^e &\sim N(0, (\sigma_\eta^e)^2), \quad z_{i0t}^e \sim N(0, (\sigma_{z_0}^e)^2) \end{aligned}$$

I map the discrete-time persistence to its continuous-time mean-reversion counterpart κ_e and the innovation variance to the OU diffusion σ_e by matching the autocovariance structure; Appendix D derives the mapping formulas. Table 4 reports the discrete-time estimates and their continuous-time images. The mean-reversion rates are similar across education groups ($\kappa_{HS} = 0.049$, $\kappa_C = 0.035$), and the diffusions are nearly identical ($\sigma_{HS} = 0.134$, $\sigma_C = 0.133$). Because college earnings are more persistent, the implied stationary standard deviation of the persistent shock is in fact slightly *higher* for college graduates (0.51 versus 0.43 for high school), so the model’s transfer-region contraction is driven by the income *level*, not by a reduction in income risk (Section 2).¹³

¹²AGMV estimate the income process separately by gender; I use their male estimates throughout, as the model abstracts from gender differences.

¹³Figure D.1 in Appendix D illustrates the implied life-cycle income profiles by education, including the deterministic age-income component, expected income for a median-ability worker, the role of OU shock volatility, and the complementarity between ability and education.

Table 4. Income Process and Age Profile

	High-School	College
<i>Ability gradient</i>		
λ_e	0.517	0.797
<i>Deterministic age profile</i>		
$\gamma_{1,e}$	0.0673	0.1197
$\gamma_{2,e} \times 1000$	-0.670	-1.191
<i>Persistent shock — discrete-time AR(1)</i>		
ρ_z^{ann}	0.952	0.966
$\sigma_{\eta,\text{ann}}^2$	0.017	0.017
$\sigma_{z_0}^2$ (entry var.)	0.059	0.094
<i>Persistent shock — continuous-time OU (mapped, fixed)</i>		
κ_e	0.049	0.035
σ_e	0.134	0.133
σ_{z_0} (entry s.d.)	0.243	0.307
stationary s.d. $\sigma_e/\sqrt{2\kappa_e}$	0.43	0.51
<i>Relative price of college labor (calibrated)</i>		
ω_C	—	0.41

Notes: Unless stated otherwise, all parameters are taken directly from [Abbott et al. \(2019\)](#) (“AGMV”, male estimates): the ability gradient λ_e , the deterministic age profile $(\gamma_{1,e}, \gamma_{2,e})$, and the discrete-time persistent-shock parameters $(\rho_z^{\text{ann}}, \sigma_{\eta}^{\text{ann}}, \sigma_{z_0})$. The continuous-time mean-reversion rates κ_e and OU diffusions σ_e are mapped from the AGMV persistence ρ_z^{ann} and innovation variance $\sigma_{\eta}^{\text{ann}}$ as described in the text and Appendix D; the same OU also governs the parent’s persistent productivity z_p over the working life. Departing from AGMV, the two education-specific transitory-income standard deviations $\sigma_{\varepsilon,e}$ are estimated by SMM rather than set to zero, so that the model can match the cross-sectional dispersion of earnings—particularly for high-school workers—that the persistent process alone understates. The relative price of college labor ω_C (ω is normalized to one for high school) is calibrated to a single target: it is set to the value (0.41) that targets the high-school-to-college income ratio and then held fixed during the SMM; the realized ratio under the estimated model is reported in Table 6. It is anchored this way rather than estimated jointly because, although it maps essentially one-to-one onto the income ratio, it also shifts college attendance; left free, the many attendance moments would pull it off the income-ratio target. The implied stationary standard deviation of the persistent shock, $\sigma_e/\sqrt{2\kappa_e}$, is 0.43 for high school and 0.51 for college.

4.3 Ability Gradient

The ability gradient λ_e in the human capital function $h(\theta, e) = (\theta/\bar{\theta})^{\lambda_e}$ governs the returns to cognitive ability by education level. I take the estimates of λ_e directly from [Abbott et al.](#)

(2019), who estimate them from the NLSY79 by regressing log hourly wages (net of the PSID age profile) on log AFQT scores, separately by education group. The key estimates—reported in Table 3—are $\hat{\lambda}_{HS} = 0.517$ and $\hat{\lambda}_C = 0.797$, implying that a one-log-point increase in ability raises college earnings by 80% but high-school earnings by only 52%. This complementarity between ability and education is central to the mechanism: the college premium is strongly increasing in ability, so removing parental transfers has heterogeneous effects across the ability distribution.

4.4 SMM Estimation

The remaining thirteen parameters—parent altruism η , psychic cost parameters κ_0 , κ_θ , and σ_κ , the parent-education psychic-cost shifter κ_w , the discount rate ρ , the warm-glow weight ϕ_b , the luxury-bequest shifter $\underline{a}_{\text{beq}}$, the collateral channel slope b_{slope} , the return-on-wealth gradient $r_{a,\text{grad}}$, the intergenerational earnings premium ζ_{pe} , and the two education-specific transitory-income standard deviations $\sigma_{\varepsilon,HS}$ and $\sigma_{\varepsilon,C}$ —are estimated by the simulated method of moments (SMM; McFadden, 1989; Pakes and Pollard, 1989). SMM is well suited to this setting because the model does not admit closed-form expressions for the moments of interest: the coupled HJB system and the endogenous transfer region must be solved numerically. Following Duffie and Singleton (1993) and Lee and Ingram (1991), I simulate the model at each candidate parameter vector and match the simulated moments to their data counterparts.

Let $\mathbf{x} \in \mathbb{R}^{13}$ denote the parameter vector. For each $\mathbf{x}$, I solve the model and simulate $M = 4,000$ dynasties, computing the model counterparts of 30 targeted moments $\mathbf{m}^{\text{sim}}(\mathbf{x})$. The estimator minimizes a weighted percentage-deviation loss with fixed diagonal weights:

$$\hat{\mathbf{x}} = \arg \min_{\mathbf{x}} \sum_m w_m \left[\frac{m_m^{\text{sim}}(\mathbf{x}) - m_m^{\text{data}}}{m_m^{\text{data}}} \right]^2, \quad (15)$$

where the weights $\{w_m\}$ are diagonal and fixed prior to estimation. Following Lee and Seshadri (2019), I do not weight all moments equally: I combine a group-size normalization with deliberate upweighting of the moments most informative about the model’s mechanisms, above all the sixteen college-attendance cells, the core target. The Online Appendix details the full weighting scheme.

4.4.1 Target Moments

The moments are organized into the following blocks, each chosen to identify specific parameters. The model is overidentified—thirteen parameters from 30 targeted moments—and although all moments contribute jointly, I flag the dominant source of identification for each parameter as the corresponding block is introduced.

College graduation rates (16 moments). College graduation rates by ability quartile $\times$ parent wealth quartile, from the NLSY97. These moments form the core of the estimation: they trace out how college attainment varies across the joint distribution of ability and parental resources, which is precisely the variation the model is designed to explain. The *ability* gradient, holding wealth fixed, identifies the psychic-cost parameters: the constant κ_0 governs the overall college rate, the log-ability gradient κ_θ how steeply attendance rises with ability, and the shock scale σ_κ the dispersion of college choices among low-ability types (a larger σ_κ raises attendance at the bottom by giving some individuals favorable cost draws). The *wealth* gradient identifies the collateral slope b_{slope} —a positive value relaxes borrowing constraints for children of wealthy parents—and, jointly with the transfer block below, the altruism parameter η ; the four high-wealth, low-ability cells, where the altruistic college gift among wealthy parents shows up most clearly, are the most informative about η .

Income (3 moments). The high-school-to-college mean income ratio and the two within-education cross-sectional variances of log earnings, all measured from individual gross earnings in the NLSY97. The two dispersions are matched by the education-specific transitory-shock standard deviations $\sigma_{\varepsilon,HS}$ and $\sigma_{\varepsilon,C}$, which the SMM estimates. The income ratio is instead targeted by the relative price of college labor ω_C , which—mapping nearly one-for-one onto this single moment—is calibrated to it outside the SMM rather than estimated jointly (Table 4).¹⁴

¹⁴Because each agent in the model represents a single individual—not a household—the income targets use individual gross earnings (wages and salary) from the NLSY97, not household income. This is consistent with the approach in Boar (2021) and Lee and Seshadri (2019), who also target individual labor income in single-agent dynastic models. The income ratio (0.69) is computed from the cross-section of individuals aged 25–37 to maximize sample size, while the two within-education variances of log earnings (0.45 for high school and 0.42 for college) are computed from ages 30–37. Earnings are individual gross earnings (wages and salary) deflated to 2016 dollars by the CPI-U; the sample conditions on positive earnings above \$5,000 and defines college graduates as those with a bachelor’s degree or higher.

Transfer moments (5 moments). The college-contingent inter-vivos transfer received *conditional on attendance* and the probability of a post-college inter-vivos transfer by parent wealth quartile. The support moment is constructed from [Abbott et al. \(2019\)](#) as the *in-college premium*: the difference between the average yearly inter-vivos transfer a male college graduate receives while enrolled and while not enrolled (\$8,203 against \$5,433 in year-2000 dollars), cumulated over the four college years and deflated to \$15.4k in 2016 dollars, conditional on attendance. Differencing the two enrollment states nets out the imputed co-residence rent—which [Abbott et al.](#) note makes up a large fraction of the raw transfer and which depends on the child’s age rather than enrollment—together with the baseline transfer a child receives regardless of college, isolating the cash support *triggered by enrolling*: precisely the object the model’s enrollment gift represents. Because in-college students co-reside less than their non-enrolled peers, the premium is a conservative lower bound on the college-contingent cash transfer, and its magnitude is corroborated by the inter-vivos transfers observed in the PSID ([Boar, 2021](#)). The model counterpart is the analogous difference: the in-college transfer—the enrollment gift τ_C plus the consumption-support flows over the college years—less the transfer received over an equal-length post-college window, averaged across college attendees.¹⁵ The four transfer-probability moments are the HRS extensive margin of Section 3.3: the probability that a parent aged 50 or older gives \$500 or more to a non-coresident child aged 26 or older within a two-year survey wave, by parent-wealth quartile. The model counterpart splits the post-college overlap into two-year windows and counts a window as positive if cumulated flow transfers plus any lump-sum transfer reach the same threshold; these moments discipline the dynamic-game transfers—and through them the lump-sum channel and the transfer boundary—separately from the college-support level.¹⁶ Because the arrival intensity λ_g is fixed and wealth-blind, the steep wealth gradient in the HRS transfer probabilities (5.5 percent in the bottom parent-wealth quartile to 24.8 percent

¹⁵Two definitional choices matter for identification. First, conditioning on attendance: an unconditional mean would conflate the attendance *rate*, already targeted by the sixteen college-rate moments, with the transfer *size*, taxing the transfer parameters twice for any attendance misfit. Second, differencing rather than matching the in-college *level*: the level bundles in-kind co-residence that the model’s cash transfers do not represent, so the premium compares like with like and avoids over-targeting transfer magnitudes; the parent consumption and wealth paths, the college block, and the extensive margin below provide complementary discipline.

¹⁶The HRS quartiles are computed within survey-wave cohorts; Appendix A details the sample construction for both sources.

in the top) can be matched only by the transfer boundary itself, which is what identifies the altruism parameter η : a higher η enlarges the transfer region, converting more Poisson arrivals into realized transfers. The college-attendance wealth gradient above provides complementary identification, since a higher η also raises the lifetime transfer burden parents face when children do not attend college.

Wealth (4 targeted moments). From the PSID linked sample: the aggregate wealth-to-income ratio, which disciplines the overall level of wealth accumulation; the median parent wealth; the p90/p50 ratio of parent wealth (7.36), which disciplines the right tail and the luxury-bequest motive; and the fraction of parents with net worth below \$25,000 (0.28), which disciplines the left mass of the parent wealth distribution. The wealth-to-income ratio and median parent wealth jointly identify the discount rate ρ —a lower ρ (more patient agents) generates more accumulation, raising both—while the p90/p50 ratio identifies the luxury-bequest shifter $\underline{a}_{\text{beq}}$ and the return-on-wealth gradient $r_{a,\text{grad}}$ through their distinct incidence on the upper tail: a positive $\underline{a}_{\text{beq}}$ makes bequests a luxury good (following [De Nardi, 2004](#)), so the rich bequeath a rising share of wealth and generate the fat right tail, while $r_{a,\text{grad}}$ raises top-end concentration without shifting the median. Following the standard of the continuous-time altruism literature ([Barczyk and Kredler, 2014a,b](#)) and the dynastic models of [Boar \(2021\)](#) and [Abbott et al. \(2019\)](#)—which discipline the parent (donor) wealth level and let the parent–child wealth relationship, and indeed the entire child wealth distribution, emerge as an equilibrium object—I target the parent wealth distribution only; the child wealth distribution and the parent–child wealth-transmission moments are left untargeted, and the child wealth distribution is reported as an out-of-sample check (Section [5.2](#)).

Education transmission (1 targeted moment). From the NLSY97: the intergenerational education gap, the difference between the college-attainment rate of children whose parents attended college and of children whose parents did not. This is the persistence statistic the mechanism speaks to—human capital rather than liquid wealth. Net of family resources, the gap identifies the parent-education taste shifter κ_w : the part of the gap that the gift and borrowing channels cannot themselves reproduce. The parent-education earnings premium ζ_{pe} , which raises a child’s earning capacity when the parent attended college,

is disciplined separately by the income block rather than by this moment.

Decumulation (1 moment). The parent wealth decumulation rate over ages 60–70, computed from the PSID linked sample. This moment identifies the warm-glow weight ϕ_b : a stronger motive slows drawdown in retirement, as parents derive utility from holding wealth at T . The rate of +2.5% per year—parents ages 60–70 are still accumulating—is consistent with evidence that pre-retirees keep saving for precautionary and warm-glow reasons (De Nardi, 2004; Hurd, 1990); this is the donor-wealth role the bequest motive plays in the continuous-time altruism literature, explaining why the elderly do not decumulate independent of how much reaches the child.

Table 5. SMM-Estimated Parameters

Parameter	Description	Value
<i>Preferences</i>		
ρ	Discount rate (annual); $\beta = e^{-\rho} = 0.90$	0.100
<i>Parental altruism & bequests</i>		
η	Altruism weight	0.09
ϕ_b	Warm-glow terminal wealth weight	7.02
$\underline{a}_{\text{beq}}$	Luxury-bequest shifter (De Nardi, \$000s)	72.8
<i>College psychic cost</i>		
κ_0	Psychic-cost constant	-0.3
κ_θ	Psychic-cost gradient in $\log(\theta)$	-0.0
σ_κ	Psychic-cost idiosyncratic shock S.D.	1.13
κ_w	Psychic-cost shifter, college parent ($e_p = C$)	-0.50
<i>College returns & financing</i>		
b_{slope}	Collateral channel: borrowing-limit slope in a_p	0.018
<i>Asset returns & income risk</i>		
$r_{a,\text{grad}}$	Return-on-wealth gradient (premium for the rich)	0.105
$\sigma_{\varepsilon,C}$	Transitory income S.D., college	0.124
$\sigma_{\varepsilon,HS}$	Transitory income S.D., high school	0.196
<i>Intergenerational transmission</i>		
ζ_{pe}	Parental-college premium in initial productivity z_0	1.36

Notes: The thirteen parameters estimated jointly by the simulated method of moments, minimizing the percentage-deviation distance between the 30 targeted simulated moments and their data counterparts. The bequest fraction α_{beq} , the intergenerational income-transmission coefficient ρ_{zp} , and the unsecured-borrowing wedge ι_b are fixed externally and reported in Table 3. Parameter definitions and the identification argument are given in Section 4. Externally set parameters, including the income process and the financial-aid schedule, are reported in Table 3.

5 Results

5.1 Model Fit

Table 6 reports fit to the targeted moments: college graduation by ability $\times$ parent-wealth quartile (Panel A), income (Panel B), transfers (Panel C), and the PSID wealth moments that discipline the warm-glow weight ϕ_b and the parent wealth distribution (Panel D). The model

matches the overall attainment level—40.7% against 42%—and both gradients. Averaged over wealth, attainment rises from 24% in the lowest ability quartile to 55% in the highest, a 31-point gradient against 36 in the data. The wealth profile is monotone and close: 33% to 54% across parent-wealth quartiles against 32% to 54%, reaching 68% for a top-ability child of a top-wealth parent against 74%.

The transfer side fits in step (Panel C): mean college support is \$14.7k against \$15.4k, and the transfer-probability wealth gradient tracks the data across quartiles (5%, 9%, 14%, 23% against 5%, 10%, 16%, 25%). The high-school-to-college income ratio is somewhat under-predicted (0.49 against 0.69), so the college premium is a little too large. On the wealth moments (Panel D) the model matches the median level of parent wealth (\$73k against \$85k), the share of parents holding little wealth (28% below \$25k, as in the data), and the pace of late-life accumulation (+2.6% against +2.5% per year). It under-predicts the aggregate wealth-to-income ratio (0.46 against 1.87) and the dispersion of parent wealth (a p90/p50 ratio of 4.1 against 7.4). Both misses have the same origin: the single-agent model, abstracting from housing and business wealth, does not reproduce the wealthiest households, who hold a disproportionate share of aggregate wealth and stretch the upper tail. The features that govern when parental support is triggered—the median, the left mass, and the decumulation pace—are matched.

Intergenerational persistence in education. The persistence the mechanism speaks to is in human capital, and there the model performs well. I target the intergenerational *education* gap—the difference between the college-attainment rate of children of college-educated parents and of children whose parents did not attend—which the model matches almost exactly (0.258 against 0.252 in the data).

Table 6. Targeted Moments: Data and Model

<i>Panel A: College Attainment by Parent Wealth and Child Ability (NLSY97)</i>					
Wealth Quartile \ Ability Quartile	1	2	3	4	
1	0.19 (0.19)	0.36 (0.24)	0.37 (0.33)	0.47 (0.53)	
2	0.24 (0.24)	0.35 (0.30)	0.42 (0.42)	0.48 (0.53)	
3	0.25 (0.26)	0.38 (0.39)	0.39 (0.51)	0.53 (0.63)	
4	0.33 (0.33)	0.49 (0.46)	0.57 (0.62)	0.68 (0.74)	

<i>Panel B: Income Moments</i>			<i>Panel D: Parent Wealth Moments</i>		
	Model	Data		Model	Data
High-School/College mean income ratio	0.49	0.69	Wealth-income ratio	0.46	1.87
High-School Var(log income)	0.35	0.45	Median wealth, parents (\$k)	72.7	85.4
College Var(log income)	0.49	0.42	p90/p50 ratio, parents	4.14	7.36
<i>Panel C: Transfer Moments</i>			Share below \$25k, parents	0.28	0.28
	Model	Data	Parent wealth growth rate	0.026	0.025
Mean transfer (\$k/yr)	14.7	15.4			
Pr(transfer > 0) — Wealth Q1	0.05	0.05			
Pr(transfer > 0) — Wealth Q2	0.09	0.10			
Pr(transfer > 0) — Wealth Q3	0.14	0.16			
Pr(transfer > 0) — Wealth Q4	0.23	0.25			

Notes: Model moments without parentheses; data moments in parentheses. Panel A: college graduation rates by parent wealth quartile (rows) and child ability quartile (columns), NLSY97. Panel B: income moments, NLSY97. Panel C: transfer moments; Panel D: parent wealth moments, PSID linked sample; wealth in thousands of 2016 dollars. Moment definitions, sources, and the targeted/untargeted distinction are given in Section 4; sample construction in Appendix A.

5.2 Untargeted Validation

The descriptive patterns of Section 3—the consumption release, the lower probability of any transfer, and the weakly larger conditional transfer—are untargeted in the SMM objective, which disciplines the attendance gradients, the wealth distribution, decumulation, and the

level of support (Section 4). They are therefore out-of-sample checks. Table 7 reports each against its model counterpart over the post-college overlap.

The model reproduces all three signs. The consumption release is +\$27.0k per year (+25%) against +\$2.5k (+8%) in the PSID; the transfer probability falls with college by 49% in the model against 21% in the HRS (Panel C of Table 9); and the conditional transfer is positive in both data and model, so the model reproduces the intensive margin. The model reproduces the qualitative pattern—including the option-value pattern that the consumption gain exceeds the fall in realized transfers, present in both the data and the model—but overstates the magnitude of the consumption release and the extensive-margin contraction.

Table 7. Untargeted Moments: Data and Model

Untargeted moment (effect of child’s college)	Data	Model
<i>Parent consumption release (college – HS, within parent-wealth quartile)</i>		
Level (\$k/yr)	+2.5	+27.0
Percent	+8%	+25%
<i>Transfer extensive margin: $Pr(\text{any transfer})$</i>		
College effect (%)	–21	–49
<i>Transfer intensive margin: $\mathbb{E}[\tau \mid \tau > 0]$, college – HS</i>		
Direction	+	+

Notes: Each row reports the effect of the child’s college attainment on a parent-side outcome that is *not* targeted in the SMM objective (Section 4); the data column reproduces the estimates of Section 3 and the model column is computed from $M = 4,000$ simulated dynasties over the post-college overlap, within parent-wealth quartile. The consumption release is the college–HS difference in mean parent consumption (\$k/yr and percent). The extensive margin is the percent change in the probability of any transfer; the model counterpart is the annual transfer probability, the data counterpart the HRS wave probability, so the levels are not directly comparable and only the college effect is. The intensive margin is the college–HS difference in the transfer conditional on a positive transfer; the data moment is measured per event and the model moment as a flow, so only the sign is comparable. Source: simulated model and Section 3.

Wealth distributions. Figure 3 compares the full parent and child wealth distributions with the data; the SMM targets only summary moments of parent wealth (Panel D of Table 6) and leaves the rest untargeted. The model reproduces the parent distribution—its right skew and the concentration of wealth among higher-income households—and the broad shape and tails of the child distribution, while placing somewhat more mass at low child wealth, so

children hold a little less by their late thirties than in the data.




Figure 3. Parent and Child Wealth Distributions: Model vs. Data

Notes: Cumulative distribution functions of parent wealth (left panel) and child wealth at ages 35–40 (right panel) in the model (simulated) and the data (PSID). Wealth is measured in thousands of 2016 dollars. The SMM targets only summary moments of parent wealth (Panel D of Table 6); the full distributions and the child wealth distribution are untargeted.

5.3 The Role of Parental Altruism in College Attainment

To quantify altruism’s role, I re-solve the model with $\eta = 0$: parents make no transfers, and children finance college through own earnings, loans, and grants. Removing altruism lowers attendance from 40.7% to 38% (Table 8), with the fall concentrated at the bottom of the ability distribution: the lowest-ability quartile drops from 24% to 17% (a 29% fall), while the middle and top quartiles move much less (2–6%). The effect is largest at the bottom because low-ability children are the marginal attendees—their private return to college is small, so parental support is pivotal to whether they enroll—whereas higher-ability children have a large enough return to attend on their own and are little affected. This isolates the *level* role of altruism: parental support finances enrollment, so removing it pushes marginal children out of college.

Table 8. College Attendance With and Without Parental Altruism

Ability quartile	Baseline Pr(C)	No altruism ($\eta = 0$) Pr(C)	Change (%)
Q1	0.24	0.17	-29
Q2	0.39	0.39	-2
Q3	0.44	0.43	-3
Q4	0.55	0.52	-6
Overall	0.41	0.38	-8

Notes: College attendance by child ability quartile in the baseline economy and in a counterfactual with parental altruism shut down ($\eta = 0$), so the parent makes no transfers and the child finances college through own earnings, loans, and grants. The final column reports the percentage change relative to the baseline. The model is re-solved at the estimated parameters with the altruism weight set to zero; statistics are computed from $M = 4,000$ simulated dynasties.

5.4 How College Reshapes the Transfer Region

College reshapes the transfer region by raising the child’s income *level*—through a higher ability gradient and a steeper earnings profile—rather than by lowering income risk. The higher level holds college children further from the boundary, so a given shock is less likely to pull them into the region where the parent transfers. Figure 4 shows this boundary shift directly: at the estimated parameters the college transfer boundary $\partial\mathcal{T}$ lies inside the high-school boundary, so a college child must be poorer before transfers become optimal. Table 9 and Figure 5 then quantify how the contracted region maps into life-cycle exposure and transfer flows.




Figure 4. Computed Transfer Boundary: College vs. High School

Notes: Transfer boundary $\partial\mathcal{T}$ in the (a_c, a_p) state space at a representative post-college age, for a median-ability child of a high-school parent in the estimated model. To the left of each boundary the parent transfers ($\tau^* > 0$); to the right the child is self-sufficient. The shaded area is the college child's transfer region $\mathcal{T}_C$; the dashed line is the high-school boundary, the solid line the college boundary. College shifts the boundary left—the child must be poorer before transfers are optimal—because higher expected income keeps child wealth further from the boundary.

Behavior inside and outside the transfer region. Panel A of Table 9 reports the child's own saving effort (which nets out the transfer itself), separating children who receive transfers from those who do not. Children who receive transfers are in the transfer region precisely because their resources are low, and both education groups run their own saving negative ($-\$2.5\text{k/yr}$ for high-school and $-\$3.0\text{k}$ for college children). Among children who receive no transfers the high-school child's own saving is close to balanced ($-\$0.4\text{k/yr}$), while the college child still runs it down ($-\$2.7\text{k/yr}$)—largely student-debt service, since interest on the outstanding loan balance depresses net resources early in the overlap. College children therefore receive transfers at the start just as high-school children do (Panel C); the college income premium pulls them out of the transfer region only as it accumulates.

Panel B isolates the behavioral distortion: the child's over-consumption relative to the full-commitment allocation evaluated at the same state, which holds resources fixed and so separates the Samaritan's dilemma from the level effect of the transfers themselves. Both education groups over-consume relative to the commitment benchmark, and the dollar amounts

scale with the child’s resources—larger outside the transfer region, where children hold more own wealth the planner would direct toward dynastic saving, and larger for college children, who earn more (inside the region, \$20.9k per year for college against \$14.0k for high school; outside, \$105.8k against \$50.1k). Because the distortion scales with resources, these dollar levels are not comparable across education groups: what makes college *reduce* the dilemma is not a smaller per-state distortion but reduced *exposure*—college children spend far less time in the transfer region (Panel C).

Exposure over the life cycle. Panel C reports exposure to the transfer region. College children are supported 48.2% of the college years—the in-college enrollment subsidy—and 10.5% of the post-college overlap, against 20.6% for high-school children. Panel (a) of Figure 5 traces this across parent age: the two profiles coincide early, while student debt holds the college child’s wealth down, and diverge as the earnings premium accumulates.

Table 9. Transfer Region Exposure and Savings Behavior

	HS Child	College Child
<i>Panel A: Child's own savings effort $s_c^{own} = \tilde{r}_c a_c + y_c - c_c$ (\$k/yr)</i>		
When receiving transfers	-2.49	-2.96
When not receiving transfers	-0.38	-2.71
Difference (in – outside region)	-2.11	-0.26
<i>Panel B: Over-consumption vs. full commitment $\Delta c_c = c_c^{MPE} - c_c^{FC}$ (\$k/yr)</i>		
In transfer region	13.97	20.91
Outside transfer region	50.05	105.75
<i>Panel C: Fraction of time receiving parental transfers</i>		
During college (ages 18–22)	—	0.482
After college (ages 22–48)	0.206	0.105

Notes: Statistics from simulated dynasties ($M = 4,000$). Panel A reports the child's own saving effort $s_c^{own} = \tilde{r}_c a_c + y_c - c_c$, which excludes the mechanical effect of transfers on wealth accumulation, separately for child-states inside and outside the transfer region; the difference reflects the composition of resources across the two regions. Panel B isolates the behavioral distortion, reporting the child's over-consumption relative to the full-commitment allocation evaluated at the same state during the overlap: the planner pools total household resources $A = a_p + a_c$ and splits them by the optimal rule $c_c^{FC} = \beta c_{tot}^{FC}$ with $\beta = \eta^{1/\gamma} / (1 + \eta^{1/\gamma})$, so the comparison holds resources fixed and removes only the strategic anticipation of transfers. A positive value indicates over-consumption relative to the commitment benchmark. Panel C reports the fraction of time the child receives parental transfers, separately for the in-college years (ages 18–22, the enrollment-subsidy channel) and the post-college period (ages 22–48, the insurance channel); the in-college figure applies only to college children.

Transfer flows by parent wealth. Transfer levels follow the same logic, but only at the top of the wealth distribution (Panel (b) of Figure 5). In the top quartile a high-school child receives \$1.8k per year against \$0.8k for a college child, a 54% reduction; the reduction holds across the wealth distribution—72% in the bottom quartile and 61–73% in the middle two—so pooled across quartiles college children receive 44% less. In levels the gap is concentrated at the top (a \$1.0k difference in the fourth quartile against \$0.1–\$0.4k below), because the lump-sum channel carries most intergenerational rebalancing (Section 2.3.4) and college debt service keeps early-career resources, and hence transfer need, close to those of high-school children. The reduction reflects the income-insurance logic of college: a dollar of tuition permanently raises the child's expected income and lowers the probability of entering the

transfer region, whereas a direct transfer smooths consumption without changing the child's income path.




Figure 5. College, Transfer-Region Exposure, and Post-College Transfers

Notes: Simulated dynasties ($M = 4,000$). Panel (a) plots the fraction of children receiving positive transfers ($\tau > 0$) at each parent age over the post-college overlap window—child age 22–48, parent age 46–72—by child education. Panel (b) reports average post-college transfers received by the child (excluding the in-college years), in thousands of 2016 dollars per year, by parent wealth quartile and child education.

6 Moral Hazard Decomposition

Parental altruism raises college attendance but generates the Samaritan's dilemma: anticipating transfers when their wealth is low, children under-save, imposing larger and more frequent transfer costs on parents. The effect on the college decision is theoretically ambiguous—anticipated in-college support subsidizes enrollment, while anticipated lifetime transfers insure non-college children against income risk. Which force dominates is quantitative. This section formalizes both forces, constructs counterfactual economies that isolate them, and reports the decomposition.

6.1 Full Commitment Benchmark and Two Opposing Forces

To isolate the effects of the Samaritan's dilemma, I first characterize the efficient allocation under full commitment. Suppose a benevolent planner commits at $t = 0$ to a state-contingent transfer schedule $\{\tau(a_p, a_c, z, t)\}$ for the entire overlap period. Because both agents pool

resources under a single plan, the problem reduces to a single-agent optimization over total household wealth $A \equiv a_p + a_c$:

$$\begin{aligned} \rho V^{\text{FC}}(A, z, t) = \max_{c_{\text{tot}}} \Big\{ & u_{\text{FC}}(c_{\text{tot}}) + V_A^{\text{FC}} [rA + y_p(t) + y_c(z, t) - c_{\text{tot}}] \\ & + \tfrac{1}{2} \sigma_a^2 A^2 V_{AA}^{\text{FC}} + \mathcal{L}_z V^{\text{FC}} + V_t^{\text{FC}} \Big\} \end{aligned} \quad (16)$$

Under CRRA preferences, the efficient allocation equates weighted marginal utilities, $u'(c_p) = \eta u'(c_c)$, so consumption shares are time-invariant: $c_c = \eta^{1/\gamma} c_p$. Total consumption $c_{\text{tot}} = c_p(1 + \eta^{1/\gamma})$ yields an aggregated flow utility $u_{\text{FC}}(c_{\text{tot}}) = \frac{c_{\text{tot}}^{1-\gamma}}{1-\gamma} (1 + \eta^{1/\gamma})^\gamma$, and the planner's HJB reduces to a single-state-variable problem. Full commitment differs from the baseline economy in three ways (Barczyk and Kredler, 2014a,b): it eliminates overconsumption by both agents, it makes the timing of transfers irrelevant—the planner pools the household's resources and fixes the consumption split, so it no longer matters when wealth moves from parent to child—and, most important for the college decision, it provides insurance without moral hazard. The child's college choice under full commitment follows the same probit structure as the baseline (equation 13), with the planner's continuation values replacing the baseline values.

Relative to full commitment, the Samaritan's dilemma creates two opposing forces on the child's college decision. From the parent's perspective, a college-educated child is cheaper to support: college permanently raises the income level, shifting the child away from the transfer boundary. Under full commitment, the parent could internalize this motive directly by conditioning transfers on e_c . In the baseline economy, however, the parent cannot commit: the transfer policy responds only to the current state (a_p, a_c, z) , and the parent transfers whenever the child is sufficiently poor, regardless of the child's education choice.

The child chooses education to maximize own continuation utility, taking the parent's equilibrium transfer policy as given. The continuation values V_C and V_{HS} embed the entire expected path of future transfers. The child does not internalize the cost transfers impose on the parent, generating overconsumption relative to the efficient allocation whenever the child's state is near or inside the transfer region. Two forces emerge from this interaction:

Force 1: Anticipated transfers subsidize enrollment. The child starts college with zero assets and reduced income ($\ell_C \cdot y_c$ net of tuition), placing the initial state deep in the transfer region. The parent—bound by time consistency—optimally transfers resources that effectively subsidize enrollment costs. This implicit subsidy raises the child’s valuation of the college path relative to full commitment, pushing attendance *above* the efficient level.

Force 2: Anticipated transfers insure non-college children. Children who forgo college face lower expected income, so they spend substantially more of their adult lives in the transfer region. The parent cannot withhold this support. The anticipation of this lifetime safety net raises the child’s valuation of the high-school path—the child can under-save and over-consume, knowing that parental transfers will materialize when income falls. This insurance reduces the child’s incentive to attend college, pushing attendance *below* the efficient level.

Which force prevails is ambiguous. The two act in opposite directions, and both are embedded in the same equilibrium transfer policy—the parent’s state-contingent support simultaneously subsidizes enrollment and insures the high-school alternative—so neither operates in isolation in the baseline economy. The balance depends on the child’s ability, the parent’s wealth, and the income process: the speed with which a college child outgrows the transfer region, set against the depth of the safety net the same child draws if she forgoes college. Resolving it requires counterfactual economies that shut down the parent’s ongoing support while preserving the $t = 0$ college gift, which I construct next.

6.2 Decomposing the Enrollment Distortion: Pure Moral Hazard and Education Targeting

Section 6.1 left a quantitative question: across the joint distribution of ability and parental wealth, does the enrollment subsidy (Force 1) or the non-college insurance (Force 2) prevail? I answer it by replacing the parent’s repeated, state-contingent support with a single lump-sum transfer at $t = 0$. The parent stays fully altruistic and still chooses the up-front gift to maximize its own value plus the η -weighted value of the child (equation 17); what these economies remove is not altruism but its *repeated* exercise—the income-contingent support a parent who cannot commit would otherwise extend whenever the child’s income fell.

Following [Barczyk and Kredler \(2014b\)](#), I call this the **dynamic interaction**. Shutting it down removes both the Samaritan’s dilemma and the insurance it provides—unlike the full-commitment benchmark, which removes the dilemma but keeps the insurance.¹⁷

Two counterfactuals do the work, and the gap between them isolates the second channel.

Counterfactual A: One-time college gift. The parent chooses a college-conditional gift τ_C at $t = 0$, paid at enrollment exactly as in the baseline:

$$\begin{aligned} \max_{\tau_C \in [0, a_p]} \quad & \Pr(C \mid \tau_C, \theta) \left[V_p^{\text{alone}}(a_p - \tau_C) + \eta V^{c, \text{alone}}(\tau_C \mid C) \right] \\ & + (1 - \Pr(C \mid \tau_C, \theta)) \left[V_p^{\text{alone}}(a_p) + \eta V^{c, \text{alone}}(0 \mid HS) \right] \end{aligned} \quad (17)$$

where $V_p^{\text{alone}}(a)$ is the parent’s remaining lifetime utility with no further interaction with the child. The child’s education choice follows the same probit rule as the baseline (equation 13), with the child-alone value gain replacing the coupled values:

$$\Pr(C \mid \tau_C, \theta) = \begin{cases} \Phi \left(\frac{G_0(\tau_C, \theta) / \bar{V} - \kappa_0 - \kappa_\theta \log \theta}{\sigma_\kappa} \right) & \text{if } G_0 > 0, \\ 0 & \text{otherwise,} \end{cases} \quad (18)$$

where $G_0(\tau_C, \theta) \equiv V^{c, \text{alone}}(\tau_C \mid C) - V^{c, \text{alone}}(0 \mid HS)$ is the child-alone value gain from college given the gift τ_C . This economy removes only the *dynamic* interaction: after the enrollment date no further transfers flow, and the child solves the lifecycle problem alone. The education decision therefore reflects fundamentals—human capital returns net of psychic cost, given the gift—rather than the anticipation of state-contingent support. Comparing baseline attendance with this economy isolates the **pure moral-hazard effect**.

Counterfactual B: Education-contingent transfer menu. The parent commits at $t = 0$ to a transfer schedule that may differ by the child’s education choice, choosing a transfer

¹⁷Appendix E provides a formal derivation of the full commitment benchmark and discusses its relationship to the one-time transfer counterfactuals. The key distinction is that the FC benchmark removes moral hazard while preserving insurance, whereas the one-time transfer counterfactuals remove both.

τ_C paid if the child attends college and a transfer τ_{HS} paid if the child chooses high school:

$$\begin{aligned} \max_{(\tau_C, \tau_{HS}) \in [0, a_p]^2} & \Pr(C \mid \tau_C, \tau_{HS}, \theta) [V_p^{\text{alone}}(a_p - \tau_C) + \eta V^{c, \text{alone}}(\tau_C \mid C)] \\ & + [1 - \Pr(C \mid \tau_C, \tau_{HS}, \theta)] [V_p^{\text{alone}}(a_p - \tau_{HS}) + \eta V^{c, \text{alone}}(\tau_{HS} \mid HS)] \end{aligned} \quad (19)$$

$$\Pr(C \mid \tau_C, \tau_{HS}, \theta) = \begin{cases} \Phi\left(\frac{G_m(\tau_C, \tau_{HS}, \theta)/\bar{V} - \kappa_0 - \kappa_\theta \log \theta}{\sigma_\kappa}\right) & \text{if } G_m > 0, \\ 0 & \text{otherwise,} \end{cases} \quad (20)$$

where the value gain is the child's own, evaluated at the education-specific endowments:

$$G_m(\tau_C, \tau_{HS}, \theta) \equiv V^{c, \text{alone}}(\tau_C \mid C) - V^{c, \text{alone}}(\tau_{HS} \mid HS).$$

where $\kappa(\theta, \varepsilon) \geq 0$ is the psychic cost of college (zero for the HS alternative). This menu is the efficient one-time-transfer benchmark: it nests Counterfactual A ($\tau_{HS} = 0$) and an unconditional transfer ($\tau_C = \tau_{HS}$), so it weakly dominates both in parent welfare. Under the menu the parent can support a high-school child directly, so a low-return child need not enroll merely to capture parental resources. Comparing the menu with Counterfactual A isolates the **education-targeting effect**.

Decomposition. Define:

$$\Delta_{\text{MH}}^{\text{pure}}(\theta, a_p) = \Pr(\text{College} \mid \text{baseline}) - \Pr(\text{College} \mid \text{college gift}) \quad (21)$$

$$\Delta_{\text{targ}}(\theta, a_p) = \Pr(\text{College} \mid \text{college gift}) - \Pr(\text{College} \mid \text{menu}) \quad (22)$$

$$\Delta_{\text{MH}}^{\text{total}}(\theta, a_p) = \Delta_{\text{MH}}^{\text{pure}} + \Delta_{\text{targ}} \quad (23)$$

The pure moral-hazard effect $\Delta_{\text{MH}}^{\text{pure}}$ is *positive* when the dynamic interaction *raises* attendance (Force 1, the in-college subsidy, prevails) and *negative* when it *lowers* attendance (Force 2, the non-college insurance, prevails). The education-targeting effect Δ_{targ} is *positive* when targeting *lowers* attendance—the parent diverting low-return children out of college and supporting them as high-school graduates instead.

Table 10. Moral Hazard Decomposition of College Attendance

	Overall	Parent Wealth Quartile			
		Q1 (Low)	Q2	Q3	Q4 (High)
<i>Panel A: College attendance rate (%), all ability types</i>					
Baseline	40.7	32.9	36.9	38.9	54.2
College gift (no MH)	43.0	27.7	35.2	45.8	63.4
$\Delta_{MH}^{\text{pure}}$ (pp)	-2.3	+5.2	+1.7	-6.9	-9.2
<i>Panel B: Pure moral-hazard effect ($\Delta_{MH}^{\text{pure}}$) (pp), by ability</i>					
Low ability		+16.5	+18.1	+12.8	+16.0
Medium ability		-0.9	-8.0	-17.4	-26.7
High ability		-5.1	-6.4	-14.6	-10.1
<i>Panel C: Optimal college gift τ_C (\$000s), by ability</i>					
Low ability		0.0	3.2	9.2	23.6
Medium ability		0.0	11.0	30.5	75.2
High ability		0.0	6.4	15.1	45.9

Notes: Panel A reports average college attendance (in percent) under the baseline and under the no-moral-hazard economy with a one-time college gift, in which parent and child interact only at $t = 0$ and the child solves the lifecycle problem alone thereafter; the “Overall” column averages across wealth quartiles. $\Delta_{MH}^{pure} = \Pr(\text{College} \mid \text{baseline}) - \Pr(\text{College} \mid \text{college gift})$, in percentage points, isolates the overconsumption distortion. Panel B disaggregates Δ_{MH}^{pure} by child ability. Panel C reports the optimal college gift τ_C by ability. The education-contingent menu (τ_C, τ_{HS}) is omitted here because the optimal high-school transfer τ_{HS} never exceeds \$4k, so it differs negligibly from the college gift. Source: simulated model ($M = 4,000$ dynasties).

Aggregate. In the estimated model the dynamic interaction *lowers* aggregate attendance (Table 10): $\Delta_{MH}^{pure} = -2.3$ percentage points (40.7% baseline against 43.0% under the one-time college gift), so on average Force 2 outweighs Force 1. The education-targeting effect is negligible throughout—the optimal high-school transfer never exceeds \$4k, so the menu barely differs from the college gift—and the total distortion is essentially the pure effect ($\Delta_{MH}^{total} = -2.2$ pp). The rest of this section therefore reads the pure effect alone.

By wealth. The small aggregate hides a sign reversal across the wealth distribution. Moral hazard *raises* attendance in the bottom two quartiles (+5.2 and +1.7 pp), where the in-college subsidy finances enrollment for liquidity-constrained families, and *lowers* it in the top two

(−6.9 and −9.2 pp), where families are wealthy enough that the non-college insurance draws marginal attendees onto the high-school path.

By ability. Panel B shows the same reversal across ability. Moral hazard *raises* attendance for low-ability children in every wealth quartile (+12.8 to +18.1 pp): for them the subsidy is pivotal to affording college at all. It *lowers* attendance for medium- and high-ability children, for whom the insurance dominates—most sharply for the medium-ability, deepening with wealth from −0.9 pp in the bottom quartile to −26.7 in the top, and −5.1 to −14.6 pp for the high-ability. The distortion the lack of commitment creates therefore operates through the *pure* moral-hazard channel—the anticipation of ongoing state-contingent support—not the parent’s inability to target transfers by education.

The asymmetry is intuitive. The dynamic interaction adds income-contingent support to both education paths, but unequally: a higher-ability child who attends college leaves the transfer region quickly, so the interaction adds little to the college path, whereas the same child who forgoes college stays nearer the region and draws a meaningful safety net. The interaction therefore raises the value of *not* attending more than the value of attending, increasingly so with ability—tipping marginal higher-ability children onto the high-school path. For low-ability children the logic reverses, because the enrollment subsidy is pivotal to affording college at all. On the margin, an altruistic parent who cannot commit ends up insuring the high-school alternative for children whose own income would otherwise make them self-sufficient.

Two roles of altruism. Read against the no-altruism comparison of Section 5.3, the two exercises separate altruism’s two roles: *whether* the parent transfers at all, and *how* the transfers are delivered—once at $t = 0$, or as ongoing, state-contingent support the parent cannot commit to withhold. Shutting altruism down entirely (Table 8) lowers attendance by three points, concentrated among low-ability children—the *level* effect, which finances enrollment at the margin of affordability. Holding that level fixed but stripping transfers of their ongoing form—keeping only the one-time college gift (Table 10)—*raises* attendance by about two points on net. The lack of commitment is therefore not a uniform force: it raises attendance for low-ability children and lowers it for the medium- and high-ability, for a small

negative effect on the whole.

7 Welfare Analysis

This section quantifies the welfare consequences of the Samaritan’s dilemma—weighing the cost of moral hazard against the insurance value of ongoing transfers—and evaluates a free-college policy.

7.1 Consumption-Equivalent Variation

I measure welfare using the consumption-equivalent variation (CEV): the permanent proportional change in consumption that makes an agent indifferent between the baseline (Samaritan’s dilemma) economy and an alternative. Under CRRA utility $u(c) = c^{1-\gamma}/(1-\gamma)$, CRRA homogeneity implies that scaling all consumption flows by a factor $(1 + \Delta)$ scales lifetime utility by $(1 + \Delta)^{1-\gamma}$. Inverting gives the CEV

$$\Delta = \left(\frac{V^{\text{alt}}}{V^{\text{base}}} \right)^{\frac{1}{1-\gamma}} - 1, \quad (24)$$

where $\Delta > 0$ means the alternative is welfare-improving.

For the child, $V^c = \int_0^T u(c_c) e^{-\rho t} dt$ and equation (24) applies directly. For the parent, $V^p = \int_0^T [u(c_p) + \eta u(c_c)] e^{-\rho t} dt$ captures both own and altruistic consumption. Because u is homogeneous of degree $1 - \gamma$, scaling both c_p and c_c by $(1 + \Delta)$ scales V^p by $(1 + \Delta)^{1-\gamma}$. The parent CEV therefore measures the permanent household-wide consumption change—own plus the altruistic component—that achieves indifference.¹⁸

7.2 Welfare Across Commitment Structures

Table 11 reports CEV for three alternative economies relative to the baseline. The first two rows of Panel A correspond to the no-moral-hazard counterfactuals from Section 6: the one-

¹⁸I compute CEV at each grid point (a_p, θ) by evaluating the value functions at the $t = 0$ decision state ($z = 0$), with the college branch at the post-gift state ($a_p - \tau_C^*, a_c = \tau_C^*$) and the high-school branch at $(a_p, 0)$, integrating over the college-choice psychic-cost shock. I then average over the simulated population: each grid CEV is weighted by the model’s stationary distribution of parent wealth and child ability, using the same quartile sort as the targeted moments, so the welfare averages are comparable to the model-fit and decomposition tables rather than to a uniform grid.

time college gift eliminates ongoing transfers and pays the child only in the college state, while the education-contingent menu also lets the parent transfer in the high-school state (τ_C, τ_{HS}) . Both allocations remove the Samaritan’s dilemma by construction—the parent commits to a transfer schedule at $t = 0$ and the child solves an independent problem thereafter. In both, the parent’s own continuation value is the single-agent parent HJB of the Online Appendix, solved with the same consumption floor, bequest motive, and wealth diffusion as the baseline coupled value, so the consumption-equivalent comparison differences like with like rather than against a permanent-income approximation.

Measured by the dynastic (η -weighted) CEV—the object the parent optimizes—both one-time-transfer economies are *worse* than the baseline (college gift -12.0% , menu -11.9%), while full commitment delivers $+15.4\%$. This contrast is the central result. The one-time transfers eliminate the Samaritan’s dilemma but sever the dynamic interaction, removing the insurance repeated transfers provide against the child’s income risk; the dynasty loses heavily, and the child—no longer exposed to moral hazard but also no longer insured against bad income draws—is itself worse off (-3.4% and -4.1%). The parent bears the larger share of the loss, for a structural reason: the one-time instruments are give-only $(\tau_C, \tau_{HS} \geq 0)$, so, unlike the full-commitment planner, the parent can never reclaim resources from the child. Its only move is a lump at $t = 0$, and because the child is liquidity-constrained then—zero assets against tuition, so its marginal value of resources is high—the altruistic parent optimally gives a large *certain* gift (up to \$75k for upper-wealth families, Panel C of Table 10), where the baseline transfers small amounts state-contingently, only in bad states and never to children who turn out fine. Bundling all support into a certain up-front gift forfeits the option value of standing ready—the very margin behind the consumption release—which is why the parent’s own value falls by more than the child’s, while the child, pocketing the lump, loses only the residual income insurance. The menu is barely distinguishable from the college gift on the dynastic measure—the high-school transfers it adds are nearly welfare-neutral for the parent even as they reshape enrollment (Section 6). Full commitment removes moral hazard while *preserving* insurance and is the only economy that raises dynastic welfare; the gain accrues entirely to the parent, who gains $+31.0\%$, while the child loses sharply (-47.6%), as commitment strips the child of the ability to extract resources by under-saving. Because

$\eta = 0.09$, the dynastic measure $V_{\text{own}}^p + \eta V_{\text{own}}^c$ is dominated by the parent’s own utility, so its positive value (+15.4%) sits far below the parent’s own gain and reflects the value of commitment *to the parent*—the rents recovered by no longer being exploited—not a Pareto improvement. Reporting the parent’s own and the child’s CEV side by side (Table 11) makes this transparent: the sign of the welfare verdict on full commitment turns entirely on how the two are weighted, whereas the one-time transfers that sever the dynamic interaction lower *both* (parent −12.8% and −12.7%, child −3.4% and −4.1%) and are unfavorable under any weighting.

Table 11. Welfare Analysis: CEV (%) Relative to Baseline

	Parent (own) CEV (%)	Child CEV (%)	Dynastic CEV (%)
<i>Panel A: Commitment Structure</i>			
Baseline (Samaritan’s dilemma)	—	—	—
College gift (no MH)	-12.84	-3.41	-11.95
Education-contingent menu (no MH)	-12.74	-4.06	-11.90
Full commitment (planner)	30.99	-47.60	15.44
<i>Panel B: Free College Policy ($\phi = 0$, lump-sum tax)</i>			
Free college	-1.39	5.18	-0.69

Notes: Consumption-equivalent variation (CEV, %) relative to the baseline. The *Parent (own)* and *Child* columns use each agent’s own-consumption value; *Dynastic* is the η -weighted value $V_p^{\text{own}} + \eta V_c^{\text{own}}$ the parent optimizes; with η small it closely tracks the parent’s own value. Panel A holds tuition at its estimated value; Panel B sets $\phi = 0$, funded by a budget-balancing lump-sum tax (footnote 19). Averages are weighted by the stationary distribution of (a_p, θ) . Source: model at estimated parameters.

The parent would like to implement full commitment but cannot, precisely because of altruism. Committing to give the child one transfer and then stand back is not time-consistent: ex ante the parent would promise to transfer τ_0 once and induce the child to self-insure, but ex post, whenever the child draws a bad shock and nears the transfer region, $\eta > 0$ makes the parent strictly prefer to transfer. The child anticipates this and under-saves, so the promise unravels—the only time-consistent (Markov-perfect) policy is the baseline. The commitment allocations are benchmarks, not policies the parent can choose unilaterally. This is why college matters: a tuition payment permanently raises the child’s income and moves it away from the transfer region for life, achieving through human capital the commitment that cash

cannot.

7.3 Free-College Policy

Free college is the natural test of the mechanism, since it lowers the price of the very investment that disciplines the dilemma. Because altruistic parents already finance much of college privately, the policy’s main effect is **crowd-out**: it substitutes public funds for private transfers, so its incidence is intra-dynastic—the gain reaches the child rather than the enrollment margin—an effect familiar from the altruistic human-capital literature ([Abbott et al., 2019](#)).

I consider a policy that eliminates tuition ($\phi = 0$) and finances the revenue shortfall through a lump-sum tax T^{LS} on initial parent wealth. No financial aid is needed when tuition is zero, so $\bar{g} = 0$. The tax is set in general equilibrium to balance the government budget: it must cover tuition for the children who enroll, while enrollment itself responds to the tax through the parent’s post-tax wealth.¹⁹

I evaluate the policy against the actual economy. Both the baseline and the free-college economy are Markov-perfect equilibria of the same no-commitment game; the policy sets $\phi = 0$ and leaves the strategic transfer relationship, and the Samaritan’s dilemma, in place. Because an altruistic parent can never commit to withholding support, a commitment version of the policy is not a feasible alternative, so the relevant comparison is the actual world against the free-college world, both without commitment.

Panel B reports the results. Free college raises enrollment only modestly (40.7% to 46.8%): because altruistic parents already finance much of college privately—through the enrollment gift and ongoing support—free tuition largely substitutes public funds for private transfers rather than financing many new entrants. The policy therefore acts mainly as a transfer to the family, and its incidence is intra-dynastic: the child gains (+5.2%) while the dynasty is roughly neutral (−0.7%, despite a lump-sum tax of present value \$12.5k), because what the subsidy relieves is the parent’s private funding burden—the intergenerational crowd-out of [Abbott et al. \(2019\)](#). The gain is concentrated where resources are scarcest: the dynastic

¹⁹The per-dynasty tax solves the fixed point $T^{\text{LS}} = 4\phi \cdot \mathbb{E}[\Pr(C \mid a_p - T^{\text{LS}}, \theta)]$, where ϕ is the annual tuition, 4ϕ the four-year cost of tuition the government covers per enrolled child, and the expectation is over the (a_p, θ) distribution. Enrollment, transfers, and value functions are all evaluated at the post-tax wealth $a_p - T^{\text{LS}}$. Parents with $a_p < T^{\text{LS}}$ have post-tax wealth clamped at zero; they still benefit because their child pays no tuition.

effect is positive in the bottom parent-wealth quartile (+0.6%) and slightly negative above it (−1.1%, −1.4%, and −0.8% in the second through top quartiles), since wealthier families pay the lump-sum tax but were already financing college privately.

8 Conclusion

This paper studies how the interaction between altruistic parents and their children shapes investment in college. The parent’s support does three things at once: it finances enrollment, it insures the child against income risk, and—because parents cannot commit to withholding it—it creates a Samaritan’s dilemma that distorts the college decision. College occupies a special place in this interaction: by permanently raising the child’s income it shrinks the transfer region, disciplining the dilemma through human capital where cash transfers cannot. In matched parent-child data, parent consumption is higher when children attend college and transfers are less frequent though conditionally larger, with the consumption gain far exceeding the fall in realized transfers—the hallmark of the mechanism: college releases the precautionary resources parents hold to stand ready to support their children.

The net effect of the Samaritan’s dilemma on college attendance is theoretically ambiguous—the lack of commitment creates an enrollment subsidy (Force 1) that raises attendance and an insurance value (Force 2) that depresses it—and which force dominates, and for whom, is a quantitative question. The structural decomposition makes that question answerable: it separates the pure moral-hazard effect from the education-targeting effect and traces both across the joint distribution of ability and parental wealth, where the balance of the two forces differs systematically across families. Whether eliminating the friction improves welfare depends on what replaces it: the dynamic interaction that distorts the college margin also insures children against income risk, so severing it—as a one-time transfer or an education-contingent menu would—leaves both parent and child worse off, a dynastic loss of roughly 12% of lifetime consumption, while only full commitment, which preserves insurance, raises welfare, and the gain accrues to the parent. The altruistic parent can adopt none of these—she cannot commit to withholding help—so they are benchmarks that value the dynamic interaction rather than alternatives to it. A free-college policy works with this logic rather than against it: it leaves the dynamic interaction and the insurance it provides in place, reallocating resources

toward the child—most where families are poorest—rather than severing the relationship that insures them. The same altruism that generates the Samaritan’s dilemma also finances college and makes the human-capital channel valuable for disciplining it, so its strength and the inability to commit jointly determine how families invest in college.

The framework carries implications for the design of college subsidies and financial aid. Because tuition policy operates on a margin embedded in an ongoing family relationship, its incidence is not confined to students: lowering the price of college reallocates resources within the dynasty, and how the gains divide between parent and child depends on the transfer relationship the policy leaves in place. The welfare effects of financial aid therefore depend on how policies interact with private transfers: grants that crowd out parental contributions may be less effective at raising enrollment than those that complement family support, while policies that expand parents’ capacity to transfer without encouraging education may deepen the Samaritan’s dilemma.

References

- Abbott, Brant, Giovanni Gallipoli, Costas Meghir, and Giovanni L Violante**, “Education policy and intergenerational transfers in equilibrium,” *Journal of Political Economy*, 2019, *127* (6), 2569–2624.
- Achdou, Yves, Jiequn Han, Jean-Michel Lasry, Pierre-Louis Lions, and Benjamin Moll**, “Income and wealth distribution in macroeconomics: A continuous-time approach,” *The Review of Economic Studies*, 2022, *89* (1), 45–86.
- Altonji, Joseph G, Fumio Hayashi, and Laurence J Kotlikoff**, “Is the extended family altruistically linked? direct tests using micro data,” *The American Economic Review*, 1992, *82*, 1177–1198.
- , —, —, and —, “Parental altruism and inter vivos transfers: Theory and evidence,” *Journal of Political Economy*, 1997, *105* (6), 1121–1166.
- Amador, Manuel, Iván Werning, and George-Marios Angeletos**, “Commitment vs. flexibility,” *Econometrica*, 2006, *74* (2), 365–396.
- Ameriks, John, Joseph Briggs, Andrew Caplin, and Mi Luo**, “Inter-generational Transfers and Precautionary Saving,” 2016 Meeting Paper 1616, Society for Economic Dynamics 2016.
- Attanasio, Orazio and José-Víctor Ríos-Rull**, “Consumption smoothing in island economies: Can public insurance reduce welfare?,” *European Economic Review*, 2000, *44* (7), 1225–1258.
- , —, **Costas Meghir**, and **Corina Mommaerts**, “Insurance in Extended Family Networks,” Working Paper 21059, National Bureau of Economic Research 2018.
- Barczyk, Daniel and Matthias Kredler**, “Altruistically motivated transfers under uncertainty,” *Quantitative Economics*, 2014, *5* (3), 705–749.
- and —, “A dynamic model of altruistically-motivated transfers,” *Review of Economic Dynamics*, 2014, *17* (2), 303–328.
- and —, “Evaluating Long-Term-Care Policy Options, Taking the Family Seriously*,” *The Review of Economic Studies*, April 2018, *85* (2), 766–809.
- and —, “Blast from the Past: The Altruism Model is Richer Than You Think,” *Journal of Economic Theory*, 2021, *198*, 105364.
- , **Sean Fahle**, and **Matthias Kredler**, “Save, Spend, or Give? A Model of Housing, Family Insurance, and Savings in Old Age,” *The Review of Economic Studies*, 2023, *90* (5), 2116–2187.
- Becker, Gary S and Nigel Tomes**, “An Equilibrium Theory of the Distribution of Income and Intergenerational Mobility,” *Journal of Political Economy*, 1979, *87* (6), 1153–1189.

- and –, “Human Capital and the Rise and Fall of Families,” *Journal of Labor Economics*, 1986, 4 (3, Part 2), S1–S39.
- Belley, Philippe and Lance Lochner**, “The changing role of family income and ability in determining educational achievement,” *Journal of Human capital*, 2007, 1 (1), 37–89.
- Bernheim, B. Douglas, Andrei Shleifer, and Lawrence H. Summers**, “The Strategic Bequest Motive,” *Journal of Political Economy*, 1985, 93 (6), 1045–1076.
- Boar, Corina**, “Dynastic Precautionary Savings,” *The Review of Economic Studies*, 2021, 88 (6), 2735–2765.
- Brandsaas, Eirik E**, “Illiquid Homeownership and the Bank of Mom and Dad,” Finance and Economics Discussion Series, Board of Governors of the Federal Reserve System 2025.
- Brown, M., J. Karl Scholz, and A. Seshadri**, “A New Test of Borrowing Constraints for Education,” *The Review of Economic Studies*, April 2012, 79 (2), 511–538.
- Bruce, Neil and Michael Waldman**, “The Rotten-Kid Theorem Meets the Samaritan’s Dilemma,” *The Quarterly Journal of Economics*, 1990, 105 (1), 155–165.
- Cameron, Stephen V and James J Heckman**, “Life cycle schooling and dynamic selection bias: Models and evidence for five cohorts of American males,” *Journal of Political Economy*, 1998, 106 (2), 262–333.
- Caucutt, Elizabeth M and Lance Lochner**, “Early and late human capital investments, borrowing constraints, and the family,” *Journal of Political Economy*, 2020, 128 (3), 1065–1147.
- Chetty, Raj, Nathaniel Hendren, Patrick Kline, and Emmanuel Saez**, “Where is the Land of Opportunity? The Geography of Intergenerational Mobility in the United States,” *The Quarterly Journal of Economics*, 2014, 129 (4), 1553–1623.
- Cox, Donald**, “Motives for private income transfers,” *Journal of Political Economy*, 1987, 95 (3), 508–546.
- Cunha, Flavio and James Heckman**, “The Economics of Human Development: The Technology of Skill Formation,” *American Economic Review*, 2007, 97 (2), 31–47.
- , –, and **Salvador Navarro**, “Separating uncertainty from heterogeneity in life cycle earnings,” *oxford Economic papers*, 2005, 57 (2), 191–261.
- , **James J. Heckman**, and **Susanne M. Schennach**, “Estimating the Technology of Cognitive and Noncognitive Skill Formation,” *Econometrica*, 2010, 78 (3), 883–931.
- Daruich, Diego and Julian Kozłowski**, “Explaining Intergenerational Mobility: The Role of Fertility and Family Transfers,” *Review of Economic Dynamics*, 2020, 36, 220–245.
- Duffie, Darrell and Kenneth J Singleton**, “Simulated Moments Estimation of Markov Models of Asset Prices,” *Econometrica*, 1993, 61 (4), 929–952.

- Fagereng, Andreas, Luigi Guiso, Davide Malacrino, and Luigi Pistaferri**, “Heterogeneity and persistence in returns to wealth,” *Econometrica*, 2020, *88* (1), 115–170.
- Gabaix, Xavier, Jean-Michel Lasry, Pierre-Louis Lions, and Benjamin Moll**, “The dynamics of inequality,” *Econometrica*, 2016, *84* (6), 2071–2111.
- Hayashi, Fumio, Joseph Altonji, and Laurence Kotlikoff**, “Risk-Sharing between and within Families,” *Econometrica: Journal of the Econometric Society*, 1996, pp. 261–294.
- Heckman, James J and Stefano Mosso**, “The economics of human development and social mobility,” *Annu. Rev. Econ.*, 2014, *6* (1), 689–733.
- , **Lance J Lochner, and Petra E Todd**, “Earnings functions, rates of return and treatment effects: The Mincer equation and beyond,” *Handbook of the Economics of Education*, 2006, *1*, 307–458.
- Hotz, V Joseph, Emily E Wiemers, Joshua Rasmussen, and Kate Maxwell Koegel**, “The role of parental wealth and income in financing children’s college attendance and its consequences,” Technical Report, National Bureau of Economic Research 2018.
- Hurd, Michael D.**, “Research on the Elderly: Economic Status, Retirement, and Consumption and Saving,” *Journal of Economic Literature*, 1990, *28* (2), 565–637.
- Kaplan, Greg**, “Moving back home: Insurance against labor market risk,” *Journal of Political Economy*, 2012, *120* (3), 446–512.
- Keane, Michael P and Kenneth I Wolpin**, “The effect of parental transfers and borrowing constraints on educational attainment,” *International Economic Review*, 2001, *42* (4), 1051–1103.
- Kocherlakota, Narayana R.**, “Implications of efficient risk sharing without commitment,” *The Review of Economic Studies*, 1996, *63* (4), 595–609.
- Laitner, John**, “Bequests, gifts, and social security,” *The Review of Economic Studies*, 1988, *55* (2), 275–299.
- Lee, Bong-Soo and Beth Fisher Ingram**, “Simulation Estimation of Time-Series Models,” *Journal of Econometrics*, 1991, *47* (2–3), 197–205.
- Lee, Sang Yoon (Tim) and Ananth Seshadri**, “On the Intergenerational Transmission of Economic Status,” *Journal of Political Economy*, April 2019, *127* (2), 855–921.
- Ligon, Ethan, Jonathan P. Thomas, and Tim Worrall**, “Informal insurance arrangements with limited commitment: Theory and evidence from village economies,” *The Review of Economic Studies*, 2002, *69* (1), 209–244.
- Lindbeck, Assar and Jörgen W Weibull**, “Altruism and time consistency: The economics of fait accompli,” *Journal of Political Economy*, 1988, *96* (6), 1165–1182.
- Lochner, Lance and Alexander Monge-Naranjo**, “The Nature of Credit Constraints and Human Capital,” *American Economic Review*, 2011, *101* (6), 2487–2529.

- Lockwood, Lee M.**, “Incidental Bequests and the Choice to Self-Insure Late-Life Risks,” *American Economic Review*, September 2018, *108* (9), 2513–2550.
- Love, David A., Michael G. Palumbo, and Paul A. Smith**, “The Trajectory of Wealth in Retirement,” *Journal of Public Economics*, 2009, *93* (1-2), 191–208.
- McFadden, Daniel**, “A Method of Simulated Moments for Estimation of Discrete Response Models Without Numerical Integration,” *Econometrica*, 1989, *57* (5), 995–1026.
- McGarry, Kathleen**, “Inter vivos transfers and intended bequests,” *Journal of Public Economics*, 1999, *73* (3), 321–351.
- , “Dynamic aspects of family transfers,” *Journal of Public Economics*, 2016, *137*, 1–13.
- Nardi, Mariacristina De**, “Wealth Inequality and Intergenerational Links,” *The Review of Economic Studies*, 2004, *71* (3), 743–768.
- **and Fang Yang**, “Bequests and heterogeneity in retirement wealth,” *European Economic Review*, November 2014, *72*, 182–196.
- , **Eric French, and John Bailey Jones**, “Why Do the Elderly Save? The Role of Medical Expenses,” *Journal of Political Economy*, 2010, *118* (1), 39–75.
- Nishiyama, S**, “Bequests, Inter Vivos Transfers, and Wealth Distribution,” *Review of Economic Dynamics*, October 2002, *5* (4), 892–931.
- Øksendal, Bernt and Agnès Sulem**, *Applied Stochastic Control of Jump Diffusions* Universitext, 2 ed., Springer-Verlag, Berlin, 2007.
- Pakes, Ariel and David Pollard**, “Simulation and the Asymptotics of Optimization Estimators,” *Econometrica*, 1989, *57* (5), 1027–1057.
- Pham, Huyên**, “Optimal stopping of controlled jump diffusion processes: a viscosity solution approach,” *Journal of Mathematical Systems, Estimation, and Control*, 1998, *8* (1), 1–27.
- Restuccia, Diego and Carlos Urrutia**, “Intergenerational Persistence of Earnings: The Role of Early and College Education,” *American Economic Review*, November 2004, *94* (5), 1354–1378.
- Solon, Gary**, “Intergenerational Income Mobility in the United States,” *The American Economic Review*, 1992, *82* (3), 393–408.
- U.S. Department of Education, Federal Student Aid**, “Standard Repayment Plan,” <https://studentaid.gov/manage-loans/repayment/plans/standard> 2025. Accessed June 2026.

A PSID Sample Construction

This appendix describes the construction of the analysis sample from the Panel Study of Income Dynamics (PSID). The sample combines four PSID data products: the Family Interview File, the Individual File, the Transition into Adulthood Supplement, and the Family Identification Mapping System. I use the thirteen biennial waves covering 1999 through 2023. The exact survey variable codes underlying each measure are documented in the replication package.

A.1 Data Sources

Family Interview File. The core source is the PSID family-level interview file, which provides household-level information on demographics, income, wealth, employment, and housing tenure for each survey wave from 1999 to 2023. A small number of supplementary measures that are not part of the main interview module—principally liquid savings and detailed industry classifications—are drawn from a companion family-level release and merged by household and wave.

Individual File. The PSID individual file provides person-level information for every sample member, including age, relationship to the household head, position within the family unit, and—beginning in 2013—years of completed education.

Transition into Adulthood Supplement. This supplement covers young adults aged roughly 18 to 28 in the 2015 and 2017 waves. I use it to confirm college enrollment directly, complementing the education-based measures available in the main files.

Family Identification Mapping System. This product provides the intergenerational links between children and their parents. Each child record carries identifiers for up to four parent types—biological father, biological mother, adoptive father, and adoptive mother—together with a generation marker that distinguishes original 1968 sample members from their children, grandchildren, and later descendants. The target sample consists of third-generation and later children (born roughly 1975–2005, reaching college age between 1993 and 2023) matched to their parents.

The remaining construction detail is provided in the Online Appendix: the full list of measures drawn from the PSID, the six-step sample-construction pipeline, the sample restrictions, the variable definitions and key constructed variables, the consumption measure, and the merge-attrition accounting.

The construction of the NLSY97 and HRS samples follows the same principles and is detailed in the Online Appendix.

B Sample Summary Statistics

Tables [B.1](#)–[B.2](#) report summary statistics for the PSID sample, stratified by average household income group. Parents in higher income groups are older, more likely to hold a college degree, more likely to own their home, and substantially wealthier: total wealth ranges from \$103,000 in the lowest income group to \$966,000 in the highest, and total consumption rises monotonically from \$36,000 to \$73,000. Among children, the college attendance rate among 18–24 year olds rises steeply with parental income. A wealth-portfolio breakdown and a consumption decomposition by expenditure category are reported in the Online Appendix.

Table B.1. Parent Demographics and Financial Characteristics by Income Group (PSID)

	By Average HH Income Group				
	\$30–60	\$60–100	\$100–160	\$160+	All
<i>Panel A: Demographics</i>					
Age of head	45.6 (16.2)	45.2 (14.4)	46.8 (13.4)	49.1 (13.2)	46.3 (14.6)
Children in household	1.3 (1.4)	1.2 (1.3)	1.1 (1.2)	1.0 (1.1)	1.1 (1.3)
Total children	2.1 (1.3)	2.0 (1.0)	2.1 (1.0)	2.2 (1.0)	2.1 (1.1)
Head has college degree	0.07 (0.26)	0.19 (0.39)	0.42 (0.49)	0.69 (0.46)	0.29 (0.45)
White	0.44 (0.50)	0.63 (0.48)	0.77 (0.42)	0.85 (0.36)	0.64 (0.48)
Black	0.45 (0.50)	0.29 (0.45)	0.18 (0.38)	0.09 (0.28)	0.28 (0.45)
Homeowner	0.50 (0.50)	0.71 (0.45)	0.84 (0.36)	0.85 (0.36)	0.70 (0.46)
<i>Panel B: Income, Wealth, and Consumption</i>					
Household income	44,223 (21,535)	76,829 (37,964)	120,169 (51,447)	242,277 (253,591)	102,569 (121,904)
Head labor income	23,523 (19,369)	40,976 (28,162)	64,728 (42,902)	123,504 (123,970)	53,930 (64,202)
Total wealth	102,555 (230,260)	200,522 (377,059)	367,076 (536,534)	965,941 (1,052,497)	326,288 (612,283)
Total consumption	36,426 (18,374)	46,131 (20,389)	57,243 (24,489)	73,130 (33,277)	50,076 (26,252)
Observations	15 096	17 356	13 011	7863	53 326
Unique parent households	2718	2770	1833	1041	8362

Notes: Means with standard deviations in parentheses. Income groups defined by parent average real household income (2016 dollars) observed between ages 25 and 60. Panel A reports demographics; Panel B reports income, wealth, and consumption. “Children in household” is the number of children under age 18 in the parent’s PSID family unit; “Total children” is the total number of children linked to the parent household via the Family Identification Mapping System (FIMS). All monetary variables in 2016 dollars, winsorized at the 1st and 99th percentiles. A detailed wealth portfolio breakdown is reported in the Online Appendix.

Table B.2. Children Descriptive Statistics by Parent Income Group (PSID)

	By Parent’s Average HH Income Group				
	\$30–60	\$60–100	\$100–160	\$160+	All
<i>Panel A: Child Demographics</i>					
Age	17.9 (19.0)	15.9 (15.7)	16.7 (13.6)	17.4 (13.4)	16.9 (16.0)
Number of siblings in sample	1.9 (1.7)	1.6 (1.2)	1.5 (1.2)	1.6 (1.1)	1.6 (1.4)
College-age (18–24)	0.14 (0.34)	0.14 (0.35)	0.16 (0.36)	0.17 (0.38)	0.15 (0.36)
<i>Panel B: Education</i>					
Years of education (age 18+)	13.7 (11.8)	13.7 (9.9)	13.8 (7.8)	14.4 (8.1)	13.9 (9.7)
Max years of education (age 25+)	16.1 (14.8)	15.7 (11.5)	15.9 (9.9)	16.5 (8.9)	16.0 (11.9)
Currently in college	0.05 (0.21)	0.06 (0.23)	0.08 (0.27)	0.10 (0.29)	0.07 (0.25)
Ever attains post-secondary ed.	0.33 (0.47)	0.37 (0.48)	0.47 (0.50)	0.56 (0.50)	0.41 (0.49)
<i>Panel C: College Rate Among 18–24 Year Olds</i>					
College attendance rate	0.32 (0.47)	0.40 (0.49)	0.47 (0.50)	0.54 (0.50)	0.42 (0.49)
College-age observations	4308	4984	4270	2907	16 469
Total child–year observations	31 731	35 484	27 026	16 962	111 203
Unique children	5680	5560	3784	2186	17 210

Notes: Means with standard deviations in parentheses. Each observation is a child-year. Children linked to parents via FIMS. Income groups refer to parent average real household income (2016 dollars). “Years of education” reported conditional on age ≥ 18 ; “Max years of education” conditional on age ≥ 25 . “College attendance rate” computed conditional on the child being aged 18–24. “Ever attains post-secondary ed.” indicates maximum observed years of education exceeds 12. Years of education are directly observed only from 2013 onward; for earlier waves, college attendance is inferred from a proxy that combines the typical college-age window with the individual’s eventual attainment (maximum years of education observed in 2013 or later exceeding 12). Individuals not observed in any post-2012 wave cannot be assigned by this proxy and are excluded from the pre-2013 attendance measure.

C Inter-Vivos Transfer Margin Regressions (HRS)

This appendix reports the extensive- and intensive-margin regressions discussed in Section 3.3: the effect of child college attainment on the probability of a positive inter-vivos transfer (Table C.1) and on the transfer amount conditional on a positive transfer (Table C.2), both estimated with parent fixed effects.

Table C.1. Extensive Margin: Effect of College on Transfer Probability (HRS)

	Extensive Margin (Linear Probability)	
	(1)	(2)
	Pr(Parent → Child)	Pr(Parent → Child)
College (BA+)	-0.031*** (0.002)	-0.010*** (0.003)
College × Q2		-0.018*** (0.004)
College × Q3		-0.024*** (0.004)
College × Q4		-0.032*** (0.005)
Parent FE	Yes	Yes
Wave FE	Yes	Yes
Dep. Var. Mean	0.144	0.144
Observations	548,164	548,164
Within R^2	0.012	0.012

Notes: Linear probability models. Dependent variable is an indicator for any positive transfer in the HRS wave. All specifications include parent fixed effects, wave dummies, child cohort controls, and parent income quartile dummies. Standard errors clustered at parent level in parentheses. * $p < .10$, ** $p < .05$, *** $p < .01$.

Table C.2. Intensive Margin: Transfer Amount Conditional on Positive Transfer (HRS)

	Intensive Margin (Amount Transfer > 0)	
	(1)	(2)
	Parent → Child	Parent → Child
College (BA+)	557* (288)	-906* (486)
College × Q2		1864*** (607)
College × Q3		965* (542)
College × Q4		1732** (684)
Parent FE	Yes	Yes
Wave FE	Yes	Yes
Dep. Var. Mean	8161	8161
Observations	78,541	78,541
Within R^2	0.002	0.002

Notes: OLS regressions on the subsample with positive transfers. Dependent variable is the real transfer amount (2016 US\$). All specifications include parent fixed effects, wave dummies, child cohort controls, and parent income quartile dummies. Standard errors clustered at parent level in parentheses. * $p < .10$, ** $p < .05$, *** $p < .01$.

D Income Process: Discrete-to-Continuous-Time Mapping and Life-Cycle Profiles

This appendix derives the mapping from the discrete-time AR(1) income process estimated by [Abbott et al. \(2019\)](#) to the continuous-time Ornstein-Uhlenbeck (OU) specification used in the model, and presents the implied life-cycle income profiles.

D.1 Mapping Derivation

The discrete-time income shock process is an AR(1) at annual frequency:

$$z_{t+1}^e = \rho_e^{\text{ann}} z_t^e + \eta_{t+1}^e, \quad \eta_{t+1}^e \sim N(0, (\sigma_\eta^e)^2). \quad (25)$$

The continuous-time counterpart is the OU process:

$$dz^e = -\kappa_e z^e dt + \sigma_e dW, \quad (26)$$

where $\kappa_e > 0$ is the mean-reversion rate and σ_e is the diffusion coefficient.

Mean-reversion rate. The exact discretization of (26) at unit time intervals ($\Delta = 1$ year) gives the conditional mean $\mathbb{E}[z_{t+1}^e | z_t^e] = e^{-\kappa_e} z_t^e$. Matching this with the AR(1) coefficient ρ_e^{ann} yields:

$$e^{-\kappa_e} = \rho_e^{\text{ann}} \implies \kappa_e = -\ln(\rho_e^{\text{ann}}). \quad (27)$$

Diffusion coefficient. The conditional variance of the discretized OU process is:

$$\text{Var}(z_{t+1}^e | z_t^e) = \frac{\sigma_e^2}{2\kappa_e} (1 - e^{-2\kappa_e}). \quad (28)$$

The conditional variance of the AR(1) is simply $(\sigma_\eta^e)^2$. Equating the two and using $e^{-\kappa_e} = \rho_e^{\text{ann}}$:

$$\begin{aligned} (\sigma_\eta^e)^2 &= \frac{\sigma_e^2}{2\kappa_e} (1 - (\rho_e^{\text{ann}})^2) \\ \implies \sigma_e &= \sigma_\eta^e \cdot \sqrt{\frac{2\kappa_e}{1 - (\rho_e^{\text{ann}})^2}}. \end{aligned} \tag{29}$$

This mapping preserves both the autocorrelation and the stationary variance $\text{Var}(z^e) = \sigma_e^2/(2\kappa_e)$ of the process. Note that the stationary variance of the OU process equals the unconditional variance of the AR(1), $(\sigma_\eta^e)^2/(1 - (\rho_e^{\text{ann}})^2)$, confirming internal consistency.

Numerical values. Applying (27)–(29) to the male estimates from [Abbott et al. \(2019\)](#) (Table 3.2): for high-school graduates the annual persistence is $\rho_{HS}^{\text{ann}} = 0.952$ with innovation variance $(\sigma_\eta^{HS})^2 = 0.017$, giving $\kappa_{HS} = -\ln(0.952) = 0.049$ and $\sigma_{HS} = \sqrt{2\kappa_{HS}(\sigma_\eta^{HS})^2/(1 - (\rho_{HS}^{\text{ann}})^2)} = 0.134$; for college graduates $\rho_C^{\text{ann}} = 0.966$ with $(\sigma_\eta^C)^2 = 0.017$, giving $\kappa_C = 0.035$ and $\sigma_C = 0.133$. Both the mean-reversion rates and the diffusions are set externally to these AGMV-implied values rather than estimated. Because college earnings are more persistent ($\rho_C > \rho_{HS}$) while the innovation variance is common, the implied stationary variance of the persistent shock, $\sigma_e^2/(2\kappa_e) = (\sigma_\eta^e)^2/(1 - (\rho_e^{\text{ann}})^2)$, is *higher* for college graduates: 0.254 (standard deviation 0.51) versus 0.182 for high school (standard deviation 0.43). The model therefore does not generate the transfer-region contraction through a reduction in income risk; the operative channel is the income level (Section 2).

D.2 Life-Cycle Profiles




Figure D.1. Income Process: Life-Cycle Profiles by Education

Notes: Panel (a): deterministic age-income component $\gamma(t, e_c) = \gamma_{1,e} t - \gamma_{2,e} t^2$ from [Abbott et al. \(2019\)](#). Panel (b): expected annual income $y_c(t) = w \cdot (\theta/\bar{\theta})^{\lambda_e} \cdot \exp(\gamma(t, e_c))$ for median ability ($\theta = 3.9$); college workers earn $\ell_C = 56\%$ of full-time income during ages 18–22 and pay net tuition. Panel (c): ± 1 stationary standard deviation of the OU income shock ($\sigma_e/\sqrt{2\kappa_e}$) around expected income. Panel (d): college income profiles for low ($\theta = 1.5$), medium ($\theta = 3.9$), and high ($\theta = 10.2$) ability workers, illustrating the complementarity between ability and education ($\lambda_C > \lambda_{HS}$).

E Full Commitment Benchmark

This appendix provides additional details on the full commitment benchmark introduced in Section 6.1 of the main text.

Under full commitment, a benevolent planner commits at $t = 0$ to a state-contingent transfer schedule $\{\tau(a_p, a_c, z, t)\}$ for the entire overlap (parent ages 42–72), solving a single-agent problem over total household wealth $A \equiv a_p + a_c$, whose law of motion carries the same multiplicative wealth diffusion $\sigma_a A dB$ as the decentralized economy (equation 16 in the main text). With CRRA preferences, the efficient allocation equates weighted marginal utilities:

$$u'(c_p) = \eta u'(c_c), \quad \text{so that} \quad c_c = \eta^{1/\gamma} c_p \quad (30)$$

Consumption shares are time-invariant. Substituting yields:

$$u_{\text{FC}}(c_{\text{tot}}) = \frac{c_{\text{tot}}^{1-\gamma}}{1-\gamma} \left(1 + \eta^{1/\gamma}\right)^\gamma \quad (31)$$

Relationship to the one-time transfer counterfactuals. The full commitment benchmark and the one-time transfer counterfactuals (Section 6.2) remove the Samaritan’s dilemma through different mechanisms. Full commitment eliminates moral hazard while preserving state-contingent insurance: the planner adjusts transfers in response to income shocks throughout the overlap period. The one-time transfer counterfactuals eliminate both moral hazard and ongoing insurance: the child receives a lump sum at $t = 0$ and is subsequently on their own. The difference between FC and the one-time transfer economies therefore isolates the *insurance-removal effect*—the change in college attendance and welfare attributable to shutting down ongoing parental support, holding the absence of moral hazard fixed.

Comparing FC with the baseline economy yields the pure moral hazard effect holding insurance constant. As discussed in Section 6.1, this comparison reveals two opposing forces: anticipated transfers subsidize enrollment (raising college attendance above the efficient level) but also insure non-college children (lowering attendance below the efficient level). The net sign depends on the child’s position in the ability–wealth distribution.

F Equilibrium Properties

This appendix derives the key properties of the Markov-Perfect Equilibrium in the continuous-time coupled problem.

F.1 Optimality Conditions

Following [Barczyk and Kredler \(2014a\)](#), the equilibrium partitions the state space into three regions: a *self-sufficient* region, where the possibility of transfers is remote; an *overconsumption* region, where no transfer currently flows ($\tau^* = 0$) but the state is drifting toward the transfer region; and the *transfer* region, where $\tau^* > 0$. The child's first-order condition $u'(c_c^*) = V_{a_c}^c$ holds throughout. In the self-sufficient region, the child's continuous-time Euler equation takes the standard form $\dot{c}_c/c_c = (r - \rho)/\gamma$. In the overconsumption region, the child anticipates that accumulating wealth will reduce future parental support, so the effective return on saving is $r + \partial\tau^*/\partial a_c$ and the Euler equation becomes:

$$\frac{\dot{c}_c}{c_c} = \frac{1}{\gamma} \left(r - \rho + \frac{\partial\tau^*}{\partial a_c} \right), \quad \frac{\partial\tau^*}{\partial a_c} < 0. \quad (32)$$

This anticipation wedge is the Samaritan's dilemma in continuous time: it lowers the child's return on saving below r and induces over-consumption *before* any transfer flows. Inside the transfer region, by contrast, the parent equalizes the marginal values of own and child wealth ($V_{a_p}^p = V_{a_c}^p$) and acts as a family dictator, setting the child's consumption directly; the child's own Euler equation is therefore not the operative condition there ([Barczyk and Kredler, 2014a](#)).

The parent's first-order condition gives $u'(c_p^*) = V_{a_p}^p$. In the no-transfer region:

$$\frac{\dot{c}_p}{c_p} = \frac{1}{\gamma} (r_{\text{eff}}(a_p) - \rho) \quad (33)$$

In the transfer region, $V_{a_p}^p = V_{a_c}^p$: the parent equalizes the marginal values of own and child wealth and makes a positive gift $\tau > 0$. The parent funds the gift out of its own resources but continues to optimize its wealth, so its drift $\dot{a}_p = r_{\text{eff}}(a_p)a_p + y_p - c_p - \tau$ need not be zero—the parent makes a *bounded* transfer (it lifts the child only toward its own optimum) and keeps saving.

F.2 Transfer Region Characterization

The complementary-slackness condition (3) partitions the state space into:

1. **No-transfer region** $\mathcal{N} = \{(a_p, a_c, z, t) : V_{a_p}^p > V_{a_c}^p\}$: Parent and child accumulate wealth independently.
2. **Transfer region** $\mathcal{T} = \{(a_p, a_c, z, t) : V_{a_p}^p = V_{a_c}^p\}$: The parent makes a positive gift that brings the child toward the altruistic target $c_c = \eta^{1/\gamma} c_p$, bounded by the child's own optimum, $\tau^* = \max\{0, \min(c_c^{\text{opt}} - R_c, \eta^{1/\gamma} c_p - R_c)\}$ with $R_c = \tilde{r}_c a_c + y_c$. The parent funds the gift from its resources but continues to save ($\dot{a}_p = r_{\text{eff}}(a_p) a_p + y_p - c_p - \tau^*$).

The no-transfer region $\mathcal{N}$ subdivides into a *self-sufficient* subregion, where the wealth distribution is balanced and future transfers are remote, and an *overconsumption* subregion bordering $\mathcal{T}$, where no transfer currently flows but the state is drifting toward $\mathcal{T}$ and the anticipation wedge of (32) is active. The transfer boundary $\partial\mathcal{T}$ is determined endogenously. Parents transfer when the child is relatively poor (low a_c), faces negative income shocks (low z), or when the parent is wealthy (high a_p). The region shrinks as the parent approaches death ($t \rightarrow T$).

Lump-sum exercise rule. The same partition governs the lump-sum channel of Section 2.3.4. At a Poisson arrival, the parent solves (5); the first-order condition for an interior lump is

$$V_{a_c}^p(a_p - \tau_L^*, a_c + \tau_L^*, z, z_p; t) = V_{a_p}^p(a_p - \tau_L^*, a_c + \tau_L^*, z, z_p; t),$$

so $\tau_L^* > 0$ if and only if $V_{a_c}^p > V_{a_p}^p$ at the pre-transfer state—the state lies strictly inside $\mathcal{T}$ —and the optimal lump moves the dynasty onto the transfer boundary $\partial\mathcal{T}$ along the constant-total-wealth line $a_p + a_c$. In $\mathcal{N}$ and on $\partial\mathcal{T}$ the opportunity expires unused ($\tau_L^* = 0$). With bounded flow transfers, income and wealth shocks can carry the state strictly into $\mathcal{T}$ between arrivals, which is what gives the lump policy positive mass: the flow props up the child's consumption while the dynasty waits for an opportunity to rebalance the stocks. Anticipated lumps also feed the child's saving incentive: the jump term in (2) adds an expected-rebalancing component to the child's continuation value that is decreasing in own

wealth, reinforcing the anticipation wedge of (32)—the Samaritan’s dilemma extends to the lump margin.

F.3 The Samaritan’s Dilemma

Both agents over-consume relative to the commitment solution. The child’s distortion arises in the overconsumption region: anticipating that accumulating wealth would move the state out of $\mathcal{T}$ and eliminate support, the child faces a reduced effective return on saving. This marginal disincentive to save is captured by $\partial\tau^*/\partial a_c$ in (32), and it operates *before* transfers flow, not inside $\mathcal{T}$ itself, where the parent dictates the allocation.

The parent’s behavior is also shaped by the lack of commitment: unable to commit to withholding support, the parent transfers whenever the child is sufficiently constrained, and the child anticipates this. Unlike a reflecting-barrier formulation, the parent does not exhaust its wealth in $\mathcal{T}$; it makes a bounded gift—lifting the child only toward its own optimum—and continues to save. The distortion therefore operates through the size and likelihood of transfers, and through the child’s reduced self-insurance, rather than through a degenerate parent saving rule.

The continuous-time formulation offers three advantages over discrete-time alternatives. First, no transfer carries within-period commitment: the flow is revisable at every instant, and the lump-sum channel operates only at exogenous arrivals, so a discrete transfer is an action taken when an opportunity strikes rather than a promise held for the length of a period. Second, the transfer boundary provides a sharp characterization of where moral hazard operates—for both the flow and the lump margins. Third, the multiplicative wealth diffusion $\sigma_a a dB$ (Section 2) renders the value functions C^2 almost everywhere and the transfer boundary smooth, so the Markov-perfect equilibrium exists in pure strategies and is computed without the non-smoothness—the “blast from the past”—that afflicts deterministic and discrete-time altruism models (Barczyk and Kredler, 2014a, 2021); the stochastic timing of lump opportunities preserves this smoothness, where deterministic transfer dates would reintroduce anticipation kinks (Proposition 1).

F.4 College and Parental Influence

The college decision at $t = 0$ is forward-looking: the child compares lifetime values under each education path, incorporating the entire future trajectory of transfers. Altruistic parents affect the decision through two channels. First, they provide generous transfers during the college years (when child income is low), effectively subsidizing attendance—the transfer region $\mathcal{T}$ is larger when $e_c = C$ during the early years. Second, the continuation value under $e_c = C$ incorporates higher future income, which feeds back into the transfer policy: wealthier children require fewer transfers, freeing parent resources. The net effect generates a positive gradient of college enrollment in parent wealth.

F.5 Computed Transfer Policy

The schematic in the main text (Figure 2) illustrates how college contracts the transfer region, and the model-computed transfer boundary $\partial\mathcal{T}$ at the estimated parameters is reported in the main text (Figure 4). Figure F.1 completes the picture with the full transfer-flow policy τ^* across the parent–child wealth state space, for a college versus a high-school child.




Figure F.1. Computed Transfer Policy τ^* by Child Education

Notes: Optimal transfer flow τ^* (thousands of 2016 dollars per year) over the parent-wealth (a_p) and child-wealth (a_c) state space, for a high-school child (left) and a college child (right), at the estimated parameters. The region where $\tau^* > 0$ is the transfer region $\mathcal{T}$; the white region is the no-transfer region $\mathcal{N}$. The two panels use independent color scales.

G College Financing Detail

This appendix details the college financial system summarized in Section 2.7, following [Abbott et al. \(2019\)](#). Let $y_p^0(\theta, e_p)$ denote the parent’s income when the child enters college (parent age 42). The annual grant is

$$g(y_p^0) = \bar{g} \cdot \max\left(0, 1 - \frac{y_p^0}{\bar{y}_{\text{efc}}}\right),$$

where $\bar{g}$ is the maximum grant and $\bar{y}_{\text{efc}}$ the parental-income level at which aid phases out completely. Borrowing up to $\bar{b}(a_p) = \bar{b}_0 + b_{\text{slope}} a_p$ is allowed at the student rate $r + \iota_s$, implemented as $a_c \geq -\bar{b}(a_p)$ during $t \in [0, 4]$; the base $\bar{b}_0 = \$18,000$ is the federal capacity available without parental collateral, and the accumulated balance amortizes so that $a_c \geq 0$ by age 32 (the standard ten-year repayment plan). The child’s wealth dynamics during college are

$$\dot{a}_c = (r + \iota_s) a_c + \ell_C \cdot y_c(t, z) - c_c - \phi_{\text{net}}(y_p^0) + \tau, \quad a_c \geq -\bar{b}(a_p),$$

where $\phi_{\text{net}}(y_p^0) = \phi - g(y_p^0)$ is net tuition. The schedule generates heterogeneity in the effective cost of college across the parental income distribution, and—as [Abbott et al. \(2019\)](#) document—government grants can crowd out the private contributions that supplement them.

H Solution Algorithm

The model is solved numerically by an upwind finite-difference scheme with implicit time-stepping, following [Achdou et al. \(2022\)](#) and the transfer-region methods of [Barczyk and Kredler \(2014a,b\)](#). This appendix summarizes the procedure and states the regularity result that underlies the smoothness of the equilibrium; the full six-stage algorithm—grid construction, the child-alone and coupled parent–child HJB solves, the college-choice stage, the dynastic simulation, and the SMM moment computation—is detailed in the Online Appendix.

Overview. The three continuous states—parent wealth a_p , child wealth a_c , and the persistent income shock z —are discretized on exponentially spaced wealth grids and a uniform Ornstein-Uhlenbeck grid, while child ability θ and education $e \in \{\text{HS}, \text{C}\}$ index separate HJB problems. The value functions are computed by backward induction. Optimal consumption follows from the upwind first-order condition $c^* = (V_a)^{-1/\gamma}$, and at each step the implicit update solves the sparse linear system

$$\left[\left(\frac{1}{\Delta t} + \rho \right) \mathbf{I} - \mathbb{A}^n \right] \mathbf{V}^n = \mathbf{u}^n + \frac{1}{\Delta t} \mathbf{V}^{n+1}$$

by sparse LU factorization, where $\mathbb{A}^n$ is the infinitesimal generator assembled from upwind drift operators, the multiplicative wealth-diffusion blocks, and the OU operator. Because altruism is one-sided, the transfer is obtained in closed form pointwise—the bounded gift (4)—so the transfer region $\mathcal{T}$ is determined without iterating on the transfer boundary. The parent and child share the same equilibrium state transitions, so the system matrix is factorized once per step and reused for both value functions; the Poisson lump-sum channel is handled by explicit operator splitting, leaving the factorization unchanged. The estimated economy and its counterfactuals are simulated for $M = 4,000$ dynasties over overlapping generations, and the 13 free parameters are estimated by a parallel CMA-ES optimizer on the weighted percentage-deviation SMM objective. The Online Appendix provides the grid specification, the stage-by-stage HJB construction, the college-choice and simulation steps, and the full SMM weighting scheme.

Regularity and the role of stochastic timing. The exogenous, exponentially-timed arrival of lump opportunities is what keeps the lump channel from disturbing the smoothness that the wealth diffusion delivers. The following statement makes this precise.

Proposition 1 (Regularity under stochastic timing). *Fix the child’s equilibrium policy and suppose the wealth diffusion is non-degenerate ($\sigma_a > 0$) on the interior of the state space. Then the parent’s value function solving the HJB (1)—with the bounded nonlocal jump term $\lambda_g [\mathcal{M}V^p - V^p]$ and the complementary-slackness condition that characterizes the transfer region $\mathcal{T}$ —is the unique viscosity solution of the associated integro-differential variational inequality. It is continuously differentiable on the interior of the state space and twice continuously differentiable away from the transfer boundary $\partial\mathcal{T}$ (hence C^2 almost everywhere), and the optimal lump maps the state onto $\partial\mathcal{T}$. The Poisson lump channel therefore introduces no anticipation discontinuity. Restricting transfers instead to a deterministic set of dates $\{t_k\}$ would impose the interior relinking conditions $V^p(\cdot, t_k^-) = \mathcal{M}V^p(\cdot, t_k^+)$ at known instants, whose non-smoothness propagates backward in time.*

Argument. The nonlocal operator $\mathcal{M}V^p(x) = \max_{\tau_L \in [0, a_p]} V^p(a_p - \tau_L, a_c + \tau_L, z, z_p; t)$ is a bounded, zeroth-order term: it is Lipschitz in x whenever V^p is, and by the envelope theorem inherits the differentiability of V^p wherever the maximizer is interior. Adding such a bounded nonlocal term to the uniformly elliptic generator $\mathcal{L}$ —uniform ellipticity following from $\sigma_a > 0$ —keeps the problem within the class of *nondegenerate* controlled jump-diffusion problems, for which the value function is the unique viscosity solution of the integro-differential variational inequality and attains $C^{1,2}$ regularity on the continuation region (Pham, 1998); the accompanying verification theorem for controlled jump diffusions is standard (Øksendal and Sulem, 2007), and the explicit treatment of the nonlocal term follows Achdou et al. (2022). Because inter-arrival times are exponential (memoryless), no instant carries positive probability of a transfer, so V_t^p develops no kink; under deterministic dates the *certain* relinking at each t_k creates a kink that propagates backward along characteristics—the “blast from the past” of Barczyk and Kredler (2021). The statement is conditional on the opponent’s policy: regularity of each agent’s value function given the other’s policy is what the diffusion-plus-Poisson structure secures, while existence of the coupled Markov-perfect equilibrium is verified numerically (smoothness of V^p , V^c , and the policies, and invariance of

the solution to grid and time-step refinement).